Name: Mohit Kerkar Div: D15C Roll No: 23

Adv DevOps Exp 12

Aim: To create a Lambda function which will log “An Image has been added” once you add an object to a specific bucket in S3

Theory:

AWS Lambda and S3 Integration:

AWS Lambda allows you to execute code in response to various events, including those triggered by Amazon S3. When an object is added to an S3 bucket, it can trigger a Lambda function to execute, allowing for event-driven processing without managing servers.

Workflow:

1. Create an S3 Bucket:
○ First, create an S3 bucket that will store the objects. This bucket will act as the trigger source for the Lambda function.
2. Create the Lambda Function:
○ Set up a new Lambda function using AWS Lambda’s console. You can choose a runtime environment like Python, Node.js, or Java.
○ Write code that logs a message like “An Image has been added” when triggered. 3. Set Up Permissions:
○ Ensure that the Lambda function has the necessary permissions to access S3. You can do this by attaching an IAM role with policies that allow reading from the bucket and writing logs to CloudWatch.
4. Configure S3 Trigger:
○ Link the S3 bucket to the Lambda function by setting up a trigger. Specify that the function should be triggered when an object is created in the bucket (e.g., when an image is uploaded).
5. Test the Setup:
○ Upload an object (e.g., an image) to the S3 bucket to test the trigger. The Lambda function should execute and log the message “An Image has been added” in AWS CloudWatch Logs.

Step 1: S3 bucket creation

Navigate to the S3 inside services section on AWS and create a new bucket

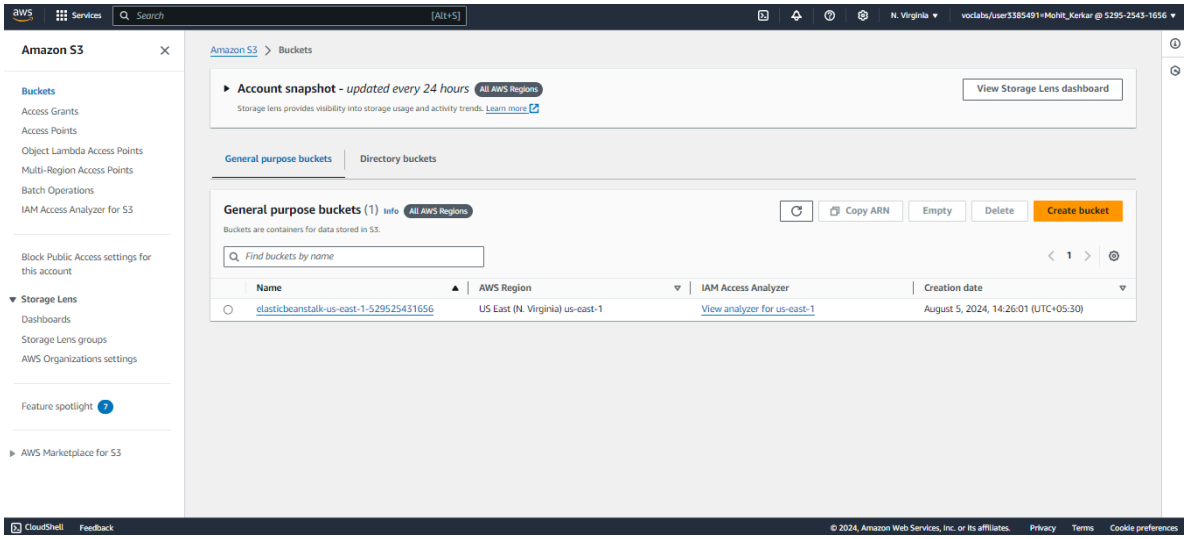

Name: Mohit Kerkar Div: D15C Roll No: 23

Give some suitable name to the bucket

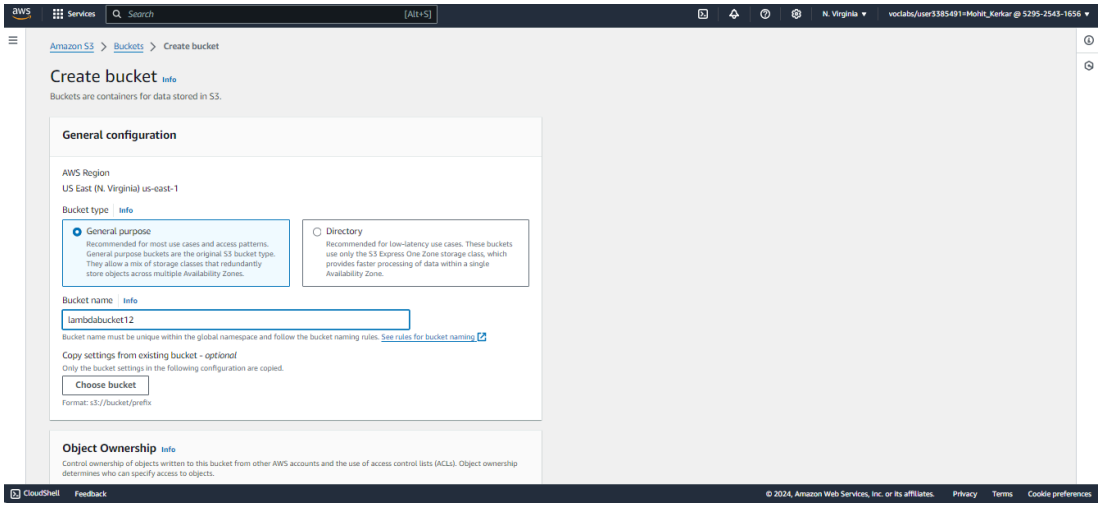

Unselect all of the default settings related to ‘Block public access’

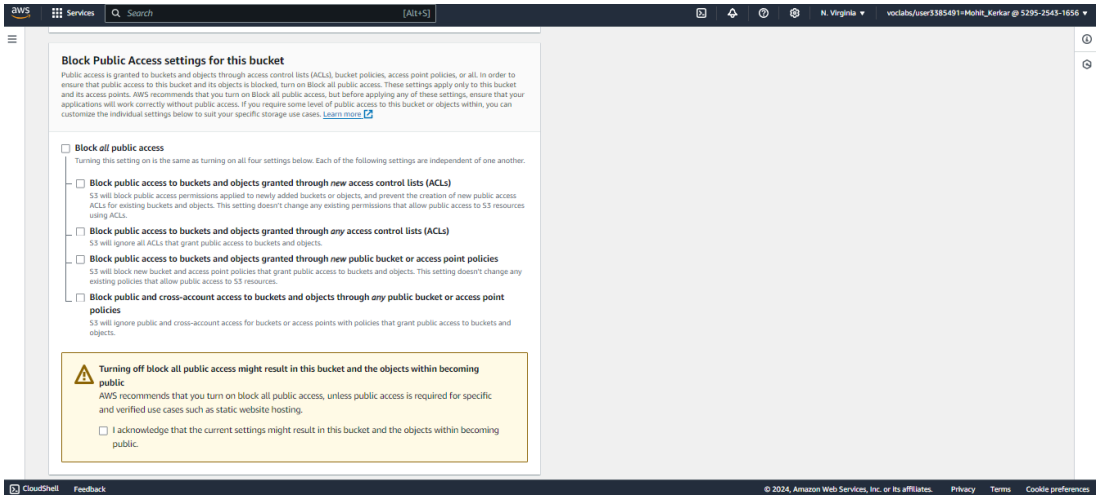

And here, you have your new S3 bucket created…

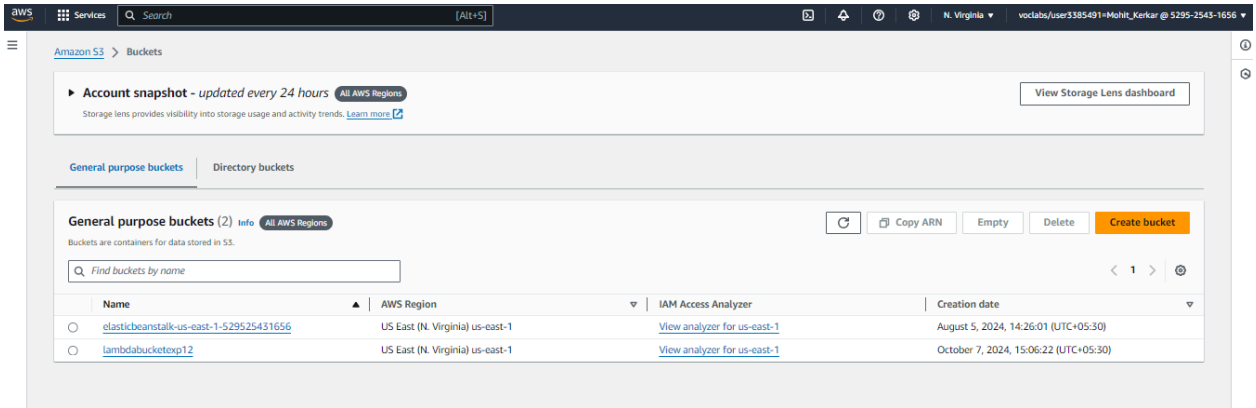

Name: Mohit Kerkar Div: D15C Roll No: 23

Step 2: Upload an image onto our bucket

Now, open the bucket and click on upload

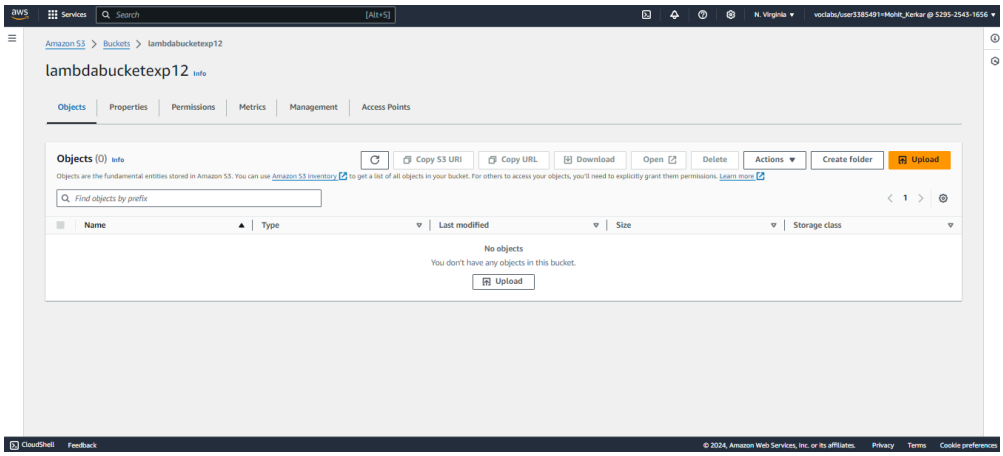

Upload a random saved image from your device onto the bucket by clicking on ‘Add files’ inside the bucket

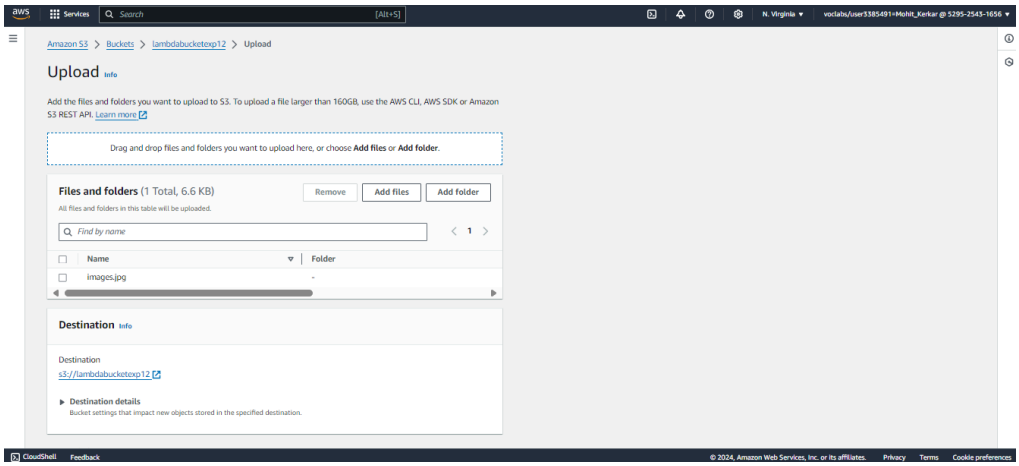

After uploading the image onto the bucket, this is the screen that will appear…

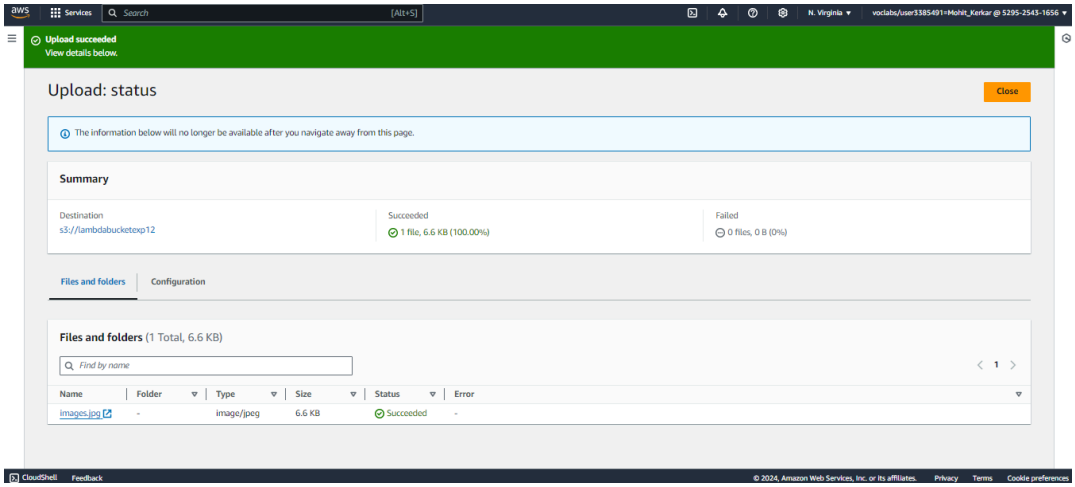

Name: Mohit Kerkar Div: D15C Roll No: 23

Step 3: Modifying Triggers inside Lambda configuration

Go back to the lambda function that we created in the previous experiment

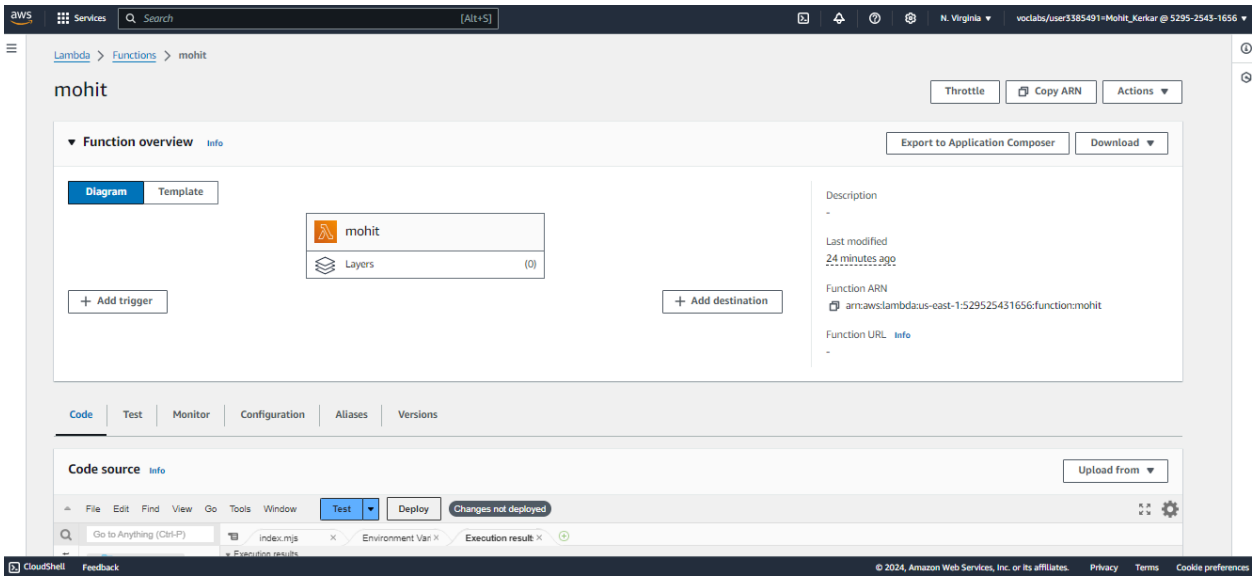

Go to the configuration section inside the function and select Triggers. This is what should appear on the screen…

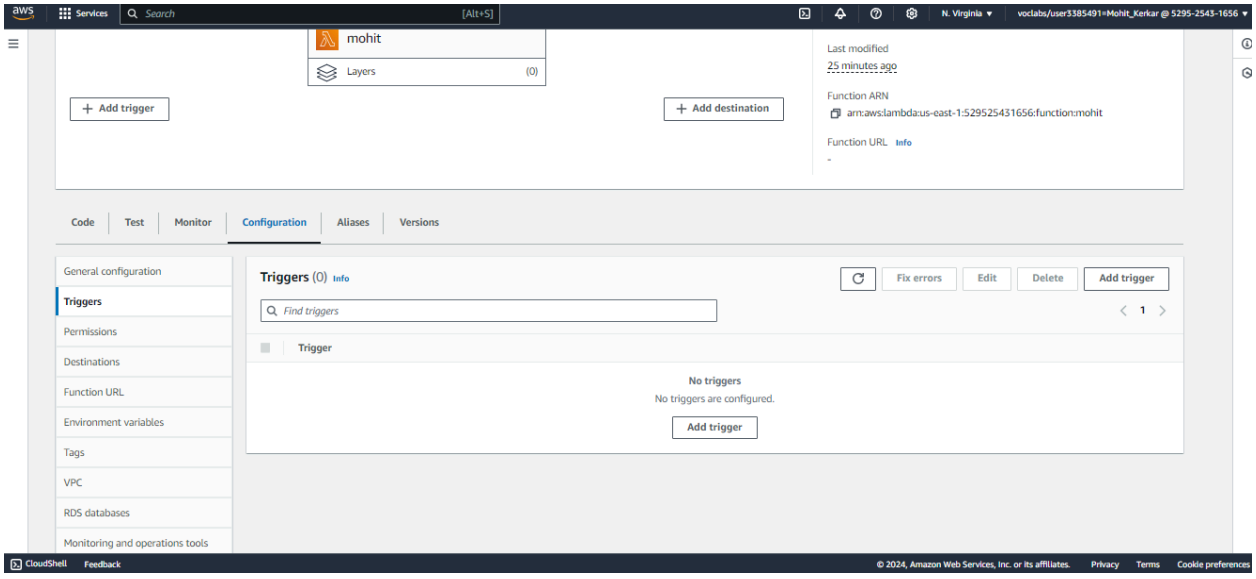

Name: Mohit Kerkar Div: D15C Roll No: 23

Select S3 bucket inside Trigger configuration to use the S3 bucket that was created earlier

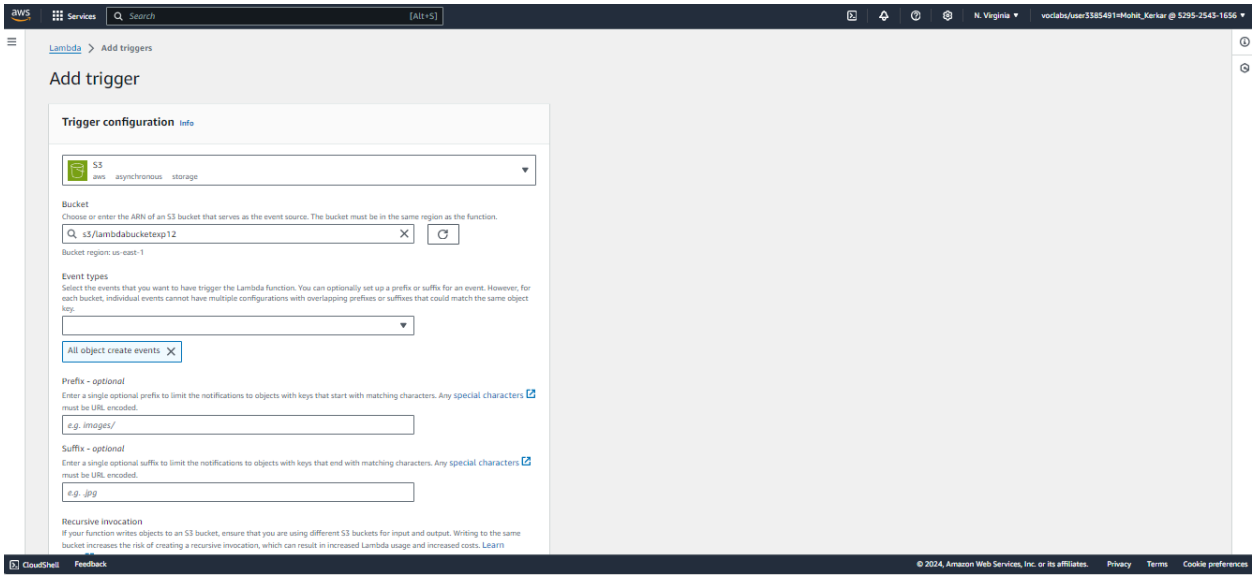

After saving the changes, we see those changes being applied in the Configuration section

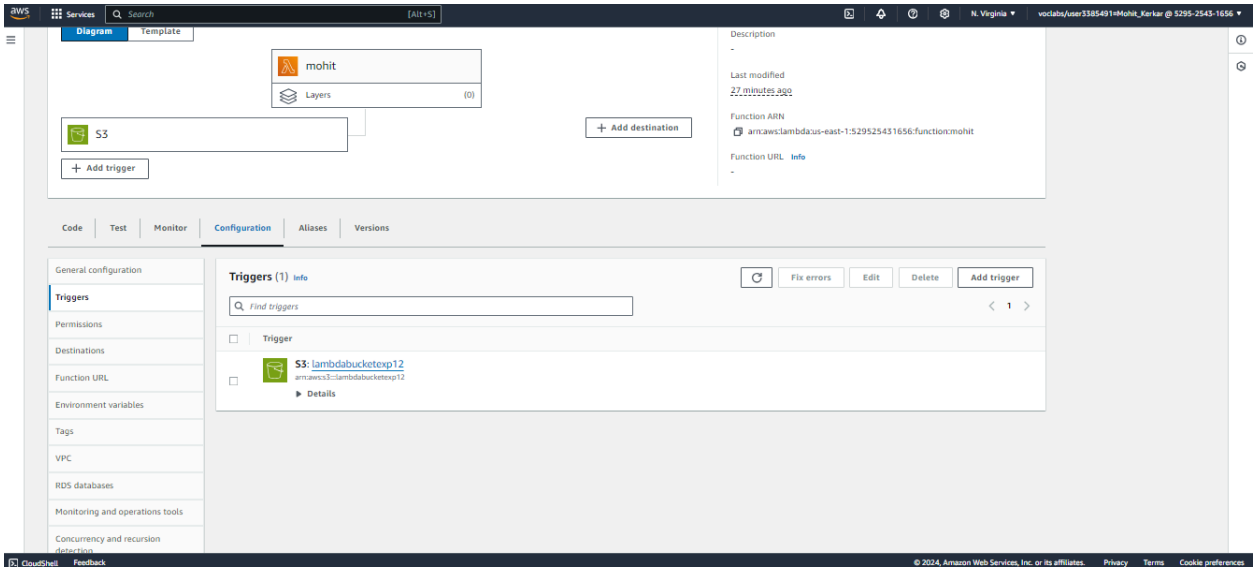

Step 4: Testing our code

In the code section, we are replace the default code with the following code:

```
export const handler = async (event) => {
  if (!event.Records || event.Records.length === 0) {
    console.error("No records found in the event.");
    return {
      statusCode: 400,
      body: JSON.stringify('No records found in the event')
    };
  }
```

```
    // Extract bucket name and object key from the event
    const record = event.Records[0];
    const bucketName = record.s3.bucket.name;
    const objectKey = decodeURIComponent(record.s3.object.key.replace(/\+/g, ' ')); // Handle
encoded keys

    console.log(`An image has been added to the bucket ${bucketName}: ${objectKey}`);
    console.log(`Event Source: ${record.eventSource}`);

    // Optionally log the entire event for debugging
    console.log('Full Event:', JSON.stringify(event, null, 2));

    return {
        statusCode: 200,
        body: JSON.stringify('Log entry created successfully!')
    };
};
```

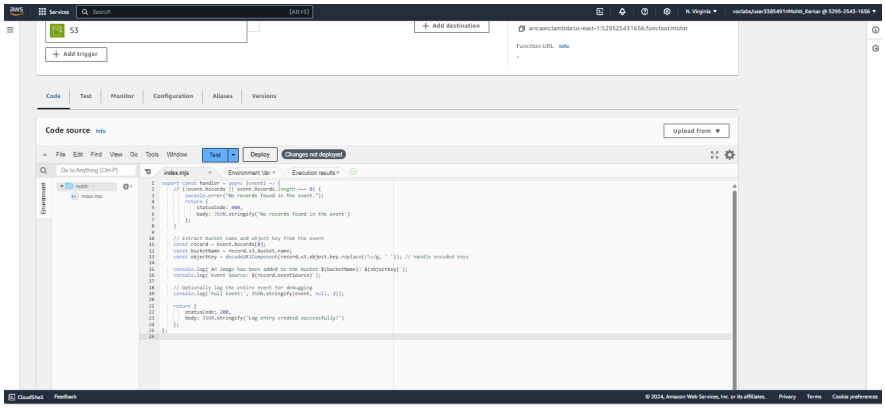

For this particular code, we are supposed to edit the configurations for our test and save the event's JSON as we see in the following screenshot

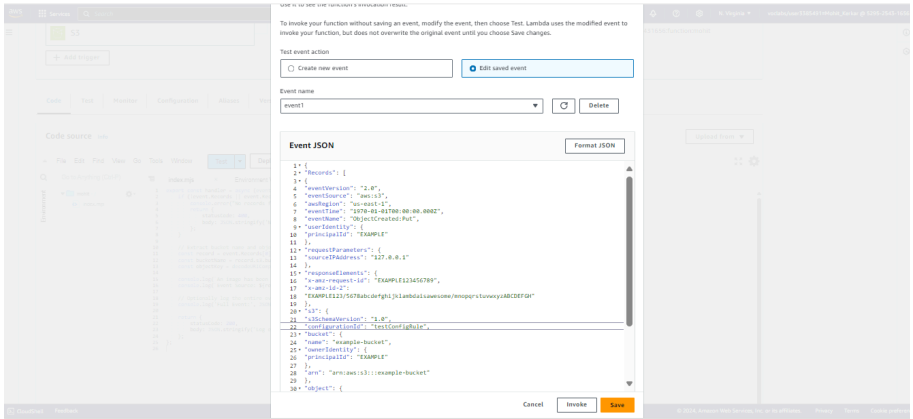

Name: Mohit Kerkar Div: D15C Roll No: 23

Click on deploy

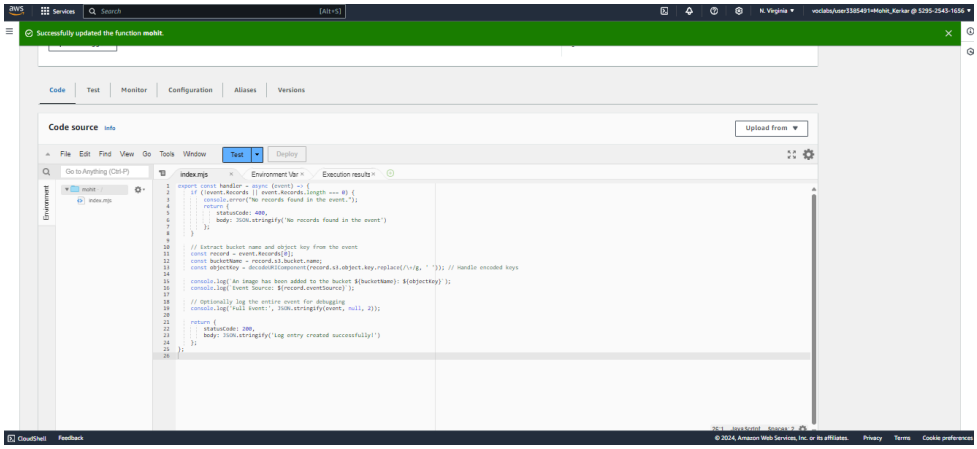

Click on test after deploying. This is the execution result that we see. Here, we see a line being mentioned inside the execution result…. ‘An image has been added to the bucket’.

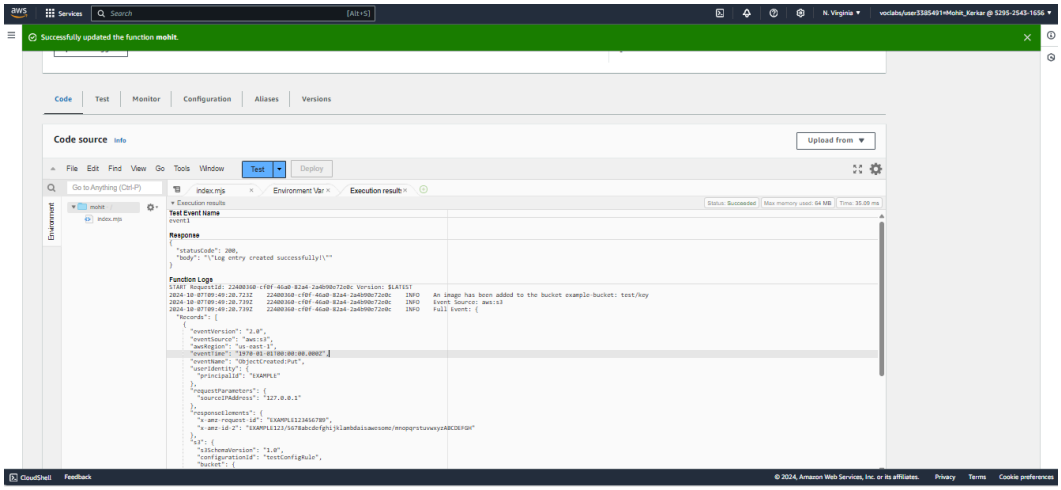

Step 5: Check logs

Search for CloudWatch. Click on logs -> logs groups

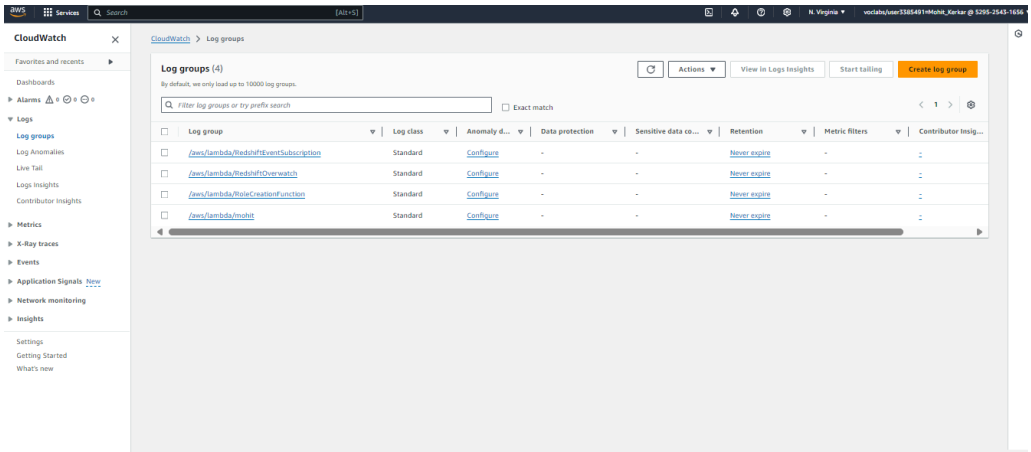

Name: Mohit Kerkar Div: D15C Roll No: 23

In here, select the logs that we want to see

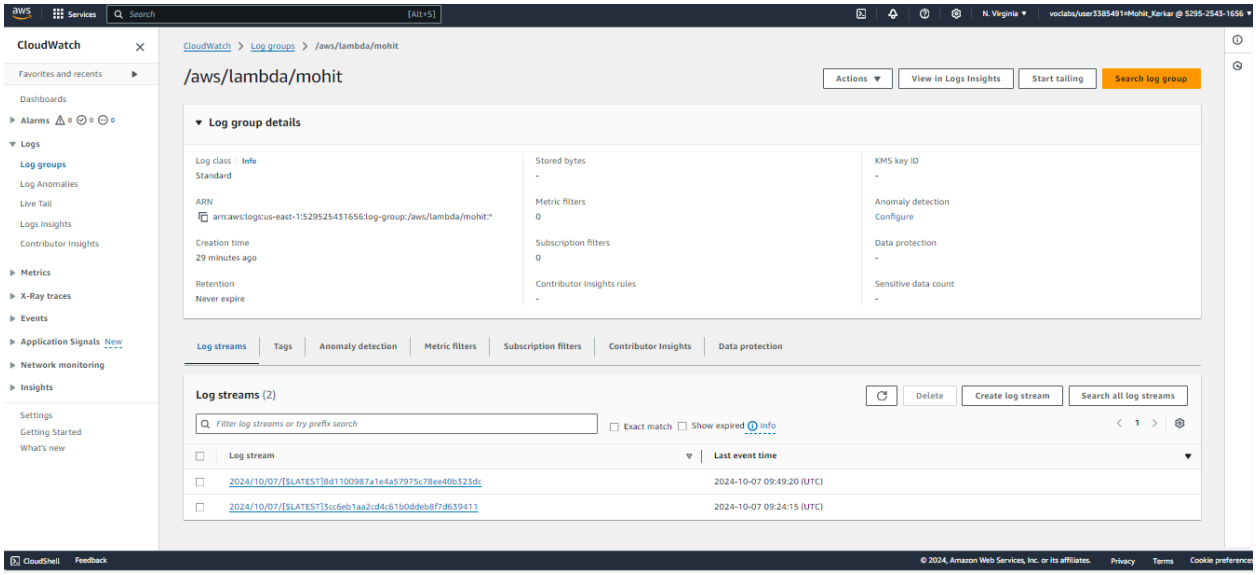

Here, we see it saying ‘An image has been added to bucket’, which matches the execution results in the test that we performed on our code

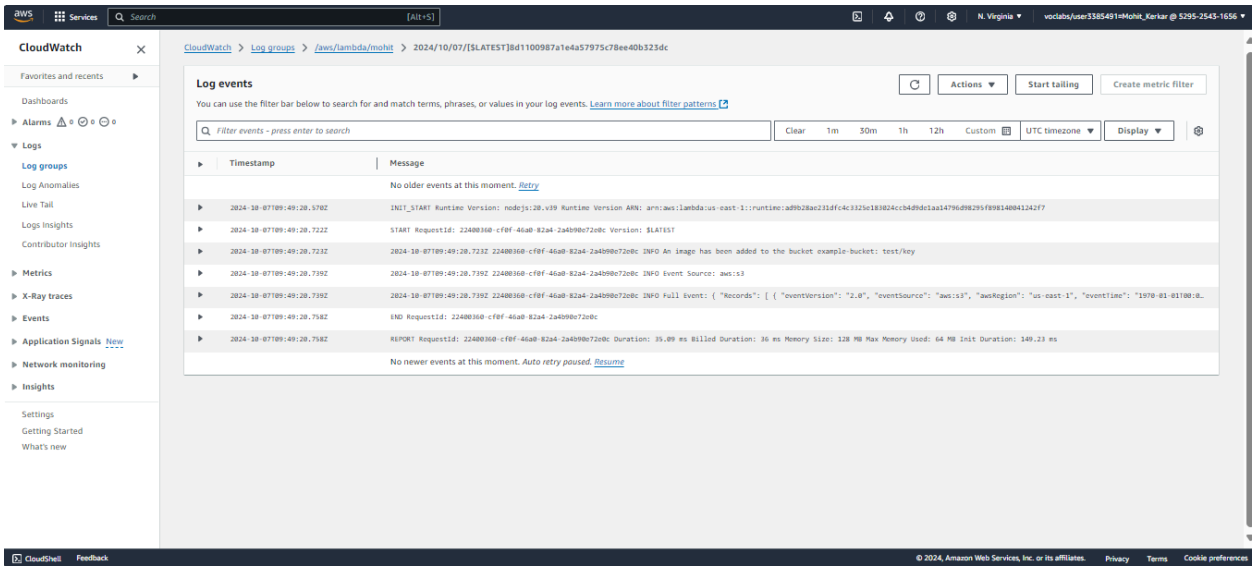

Conclusion: Integrating AWS Lambda with S3 allows for real-time, automated processing of events such as file uploads. In this example, a Lambda function is configured to log a message whenever an image is added to a specific S3 bucket. This setup demonstrates the power and flexibility of serverless computing by automating tasks without requiring manual intervention or server management. By leveraging AWS Lambda, developers can efficiently handle event-driven workflows, reduce operational overhead, and quickly deploy scalable solutions that respond to specific actions within cloud environments.