

Adv DevOps Lab Exp 03

Aim: To understand the Kubernetes Cluster Architecture, install and Spin Up a Kubernetes Cluster on Linux Machines/Cloud

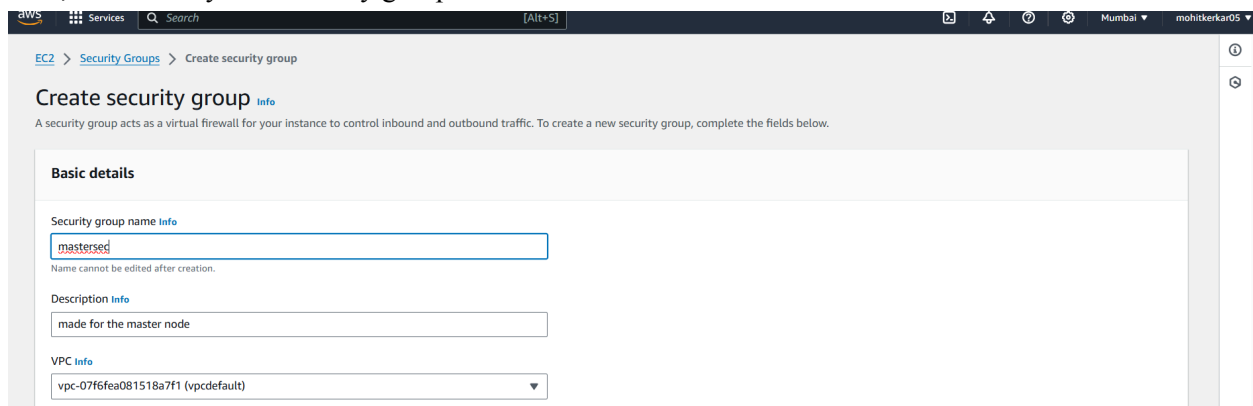
(I have performed this experiment on my personal AWS account)

Step 1: Create Key-pair, Security groups and required default VPCs

I started off by creating two separate security groups i.e one for the master instance (kubeadm to be initialized within it) and the other for the 2 worker instances. For creating any EC2 security groups, we require a VPC (Virtual private cloud) and a subnet along with it .. so that we can work with or allow inbound and outbound traffic and communication. Make sure you have a default VPC and a subnet already created which can be used for this experiment.

Also, make sure that you have a key pair installed of the type RSA with .pem extension on your local machine. Save it at a place which is accessible and where you can work from your terminal.

Here, i created my first security group

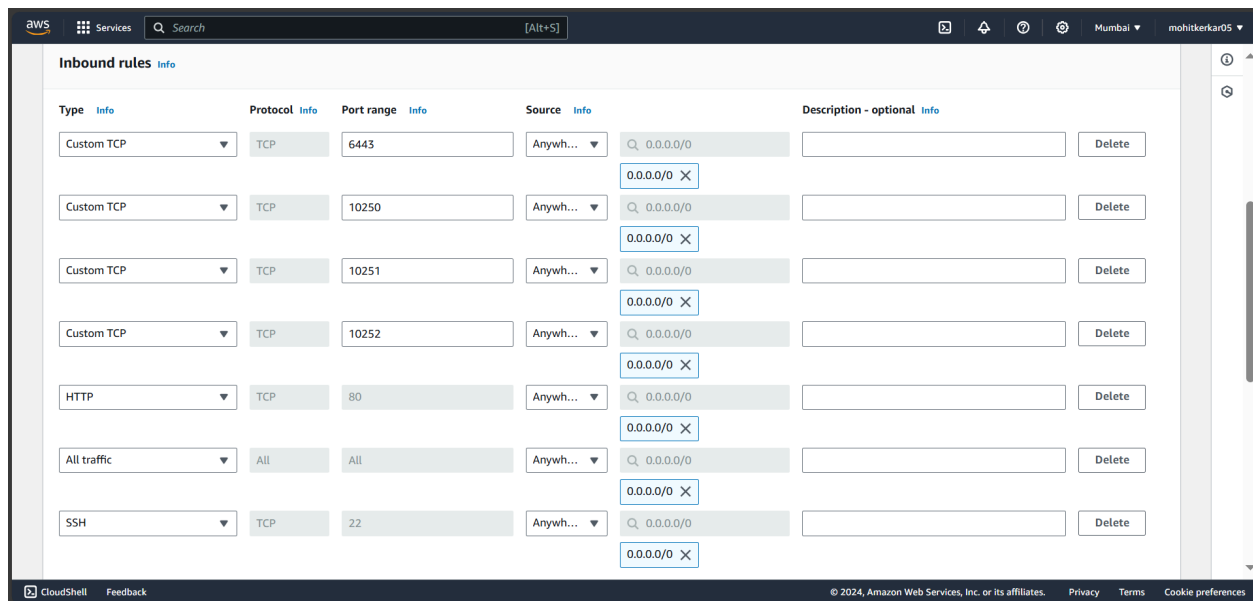


The screenshot shows the AWS Management Console interface for creating a new security group. The breadcrumb navigation indicates the path: EC2 > Security Groups > Create security group. The main heading is 'Create security group' with an 'Info' link. Below the heading, a note states: 'A security group acts as a virtual firewall for your instance to control inbound and outbound traffic. To create a new security group, complete the fields below.'

The 'Basic details' section contains the following fields:

- Security group name:** A text input field containing 'mastered'. A red warning box below it says 'Name cannot be edited after creation.'
- Description:** A text input field containing 'made for the master node'.
- VPC:** A dropdown menu showing 'vpc-07f6fea081518a7f1 (vpcdefault)'.

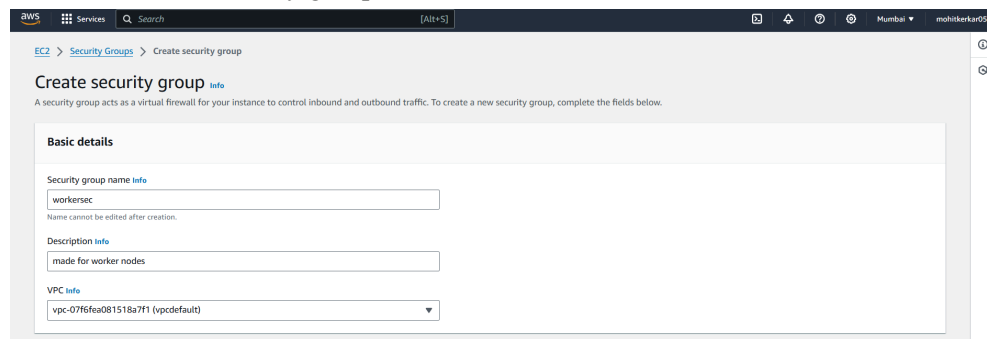
I modified the inbound rules in the following fashion for the master node



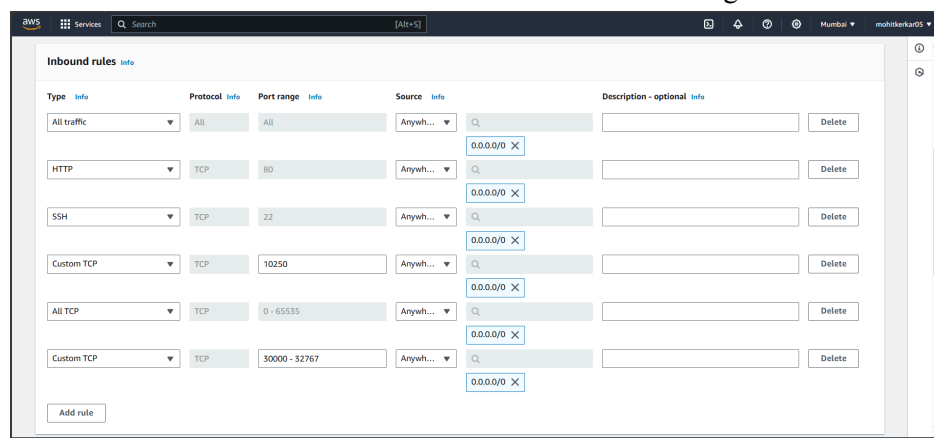
The screenshot shows the 'Inbound rules' page for the 'mastered' security group. The page has a table of rules with columns: Type, Protocol, Port range, Source, and Description - optional. Each rule has a 'Delete' button.

Type	Protocol	Port range	Source	Description - optional	Action
Custom TCP	TCP	6443	Anywh...	0.0.0.0/0	Delete
Custom TCP	TCP	10250	Anywh...	0.0.0.0/0	Delete
Custom TCP	TCP	10251	Anywh...	0.0.0.0/0	Delete
Custom TCP	TCP	10252	Anywh...	0.0.0.0/0	Delete
HTTP	TCP	80	Anywh...	0.0.0.0/0	Delete
All traffic	All	All	Anywh...	0.0.0.0/0	Delete
SSH	TCP	22	Anywh...	0.0.0.0/0	Delete

Then, i created a security group for the worker nodes



Edited the inbound rules for the worker nodes in the following fashion

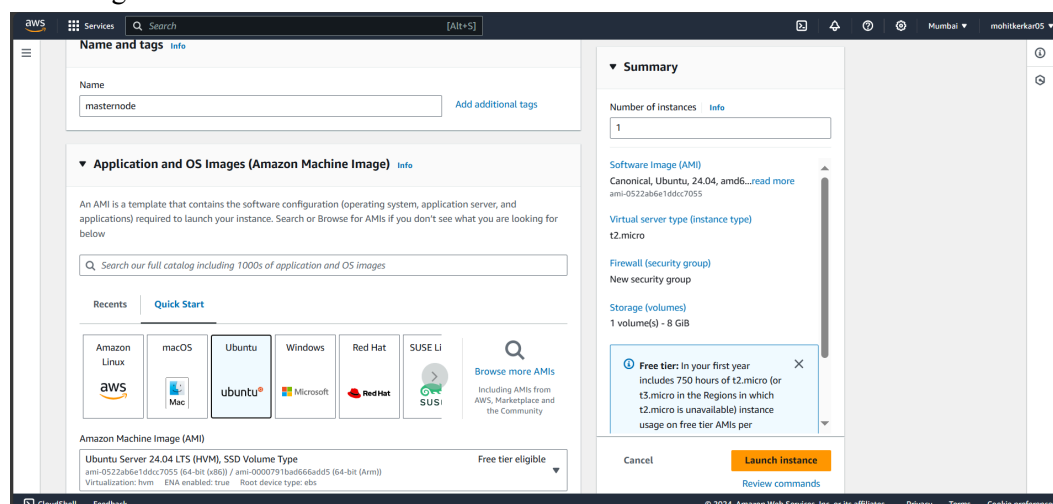


Step 2: Create 3 EC2 instances i.e one master and other 2 workers

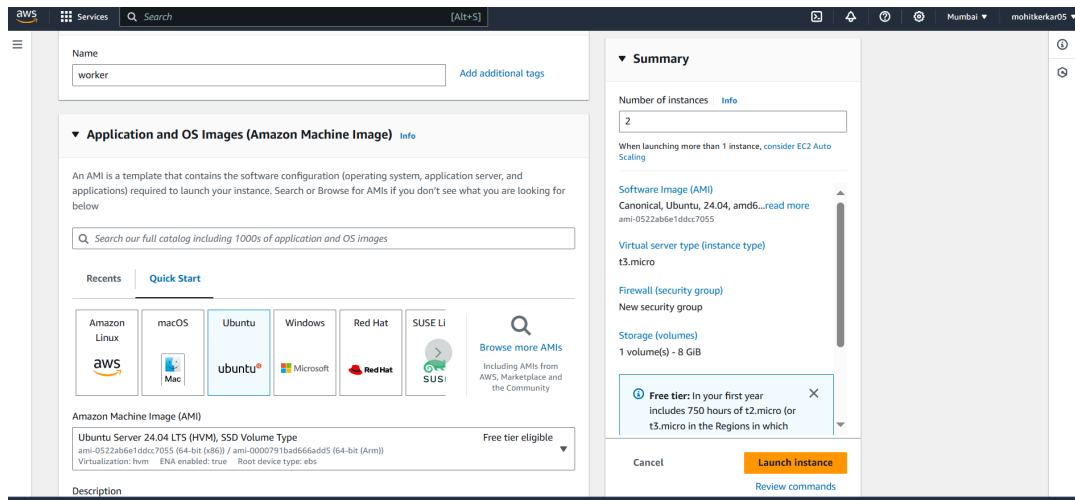
Now, in order to work with the instances, navigate to EC2 instances section and launch a few instances. Launch one instance and name it as masternode and the other two as worker.

I selected AMI as ubuntu among the list of OS that we shown available in the free tier eligibility list

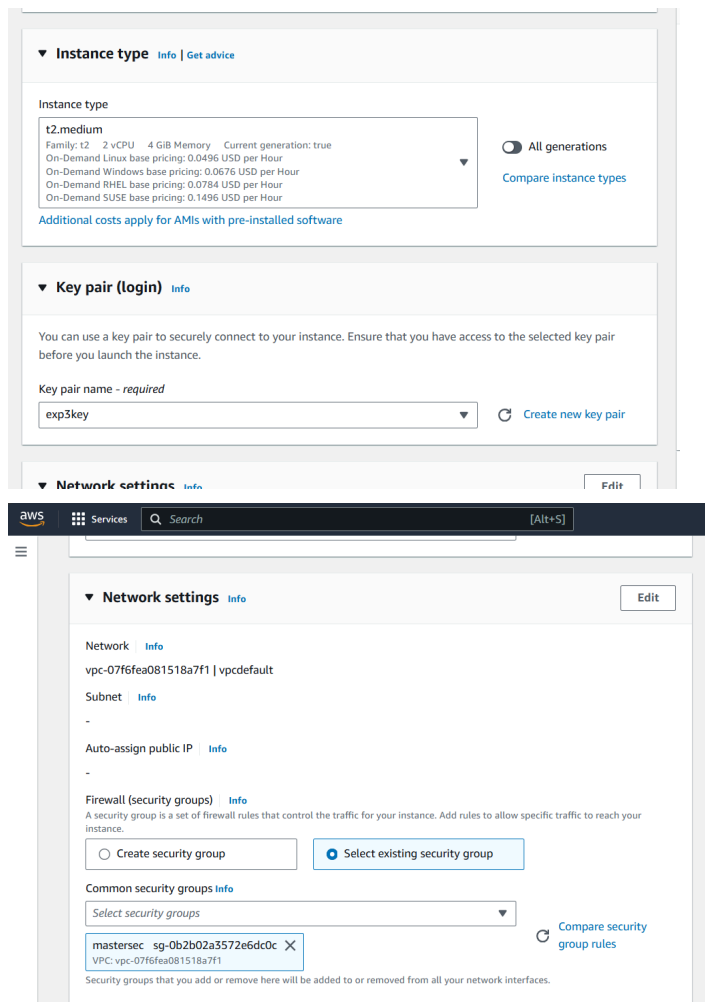
Following is the master instance node:




Following is the worker instance node:



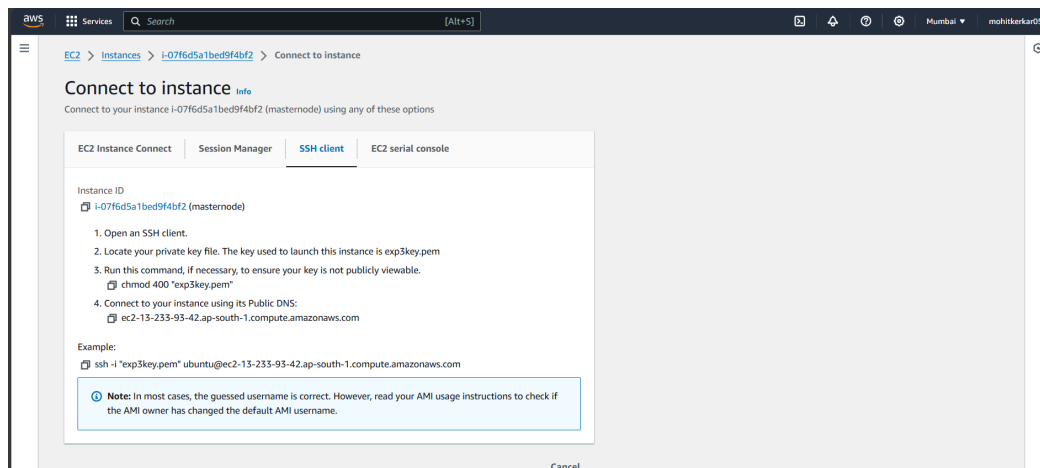
IMPORTANT: Select t2.medium as the instance type for all 3 instances, as kubernetes requires at least 2 CPUs to work with and sufficient amount of other resources as well



After launching all the 3 instances, select one instance and press 'Connect'



Navigate to SSH client section and copy the command that's listed for connecting to the node remotely `ssh -i "exp3key.pem" ubuntu@ec2-13-233-93-42.ap-south-1.compute.amazonaws.com`...for the master node. Similarly, do this for all nodes.



Now, open 3 separate terminals and run change directories to the folder which contains the key we created earlier.

Run their respective SSH commands (one in each terminal)

This helps us log onto those instances individually and remotely and work on them separately.

```
ubuntu@ip-172-31-37-198: ~
PS C:\Users\Dell> cd C:\Users\Dell\Desktop\keypair
PS C:\Users\Dell\Desktop\keypair> ssh -i "exp3key.pem" ubuntu@ec2-3-111-188-84.ap-south-1.compute.amazonaws.com

The authenticity of host 'ec2-3-111-188-84.ap-south-1.compute.amazonaws.com (3.111.188.84)' can't be established.
ECDSA key fingerprint is SHA256:Lm8ajmtdu/3LBU1quOmUjWlCHZLqFwhD1jaUCAQAoqM.
Are you sure you want to continue connecting (yes/no/[fingerprint])? yes
Warning: Permanently added 'ec2-3-111-188-84.ap-south-1.compute.amazonaws.com,3.111.188.84' (ECDSA) to the list of known hosts.
Welcome to Ubuntu 24.04 LTS (GNU/Linux 6.8.0-1012-aws x86_64)

 * Documentation:  https://help.ubuntu.com
 * Management:    https://landscape.canonical.com
 * Support:       https://ubuntu.com/pro

System information as of Thu Sep 26 09:43:59 UTC 2024

System load:  0.34               Processes:            118
Usage of /:   22.8% of 6.71GB    Users logged in:     0
Memory usage: 6%                IPv4 address for enX0: 172.31.37.198
Swap usage:   0%

Expanded Security Maintenance for Applications is not enabled.

0 updates can be applied immediately.
```

```

PS C:\Windows\system32> cd C:\Users\Dell\Desktop\keypair
PS C:\Users\Dell\Desktop\keypair> ssh -i "exp3key.pem" ubuntu@ec2-13-234-38-84.ap-south-1.compute.amazonaws.com

The authenticity of host 'ec2-13-234-38-84.ap-south-1.compute.amazonaws.com (13.234.38.84)' can't be established.
ECDSA key fingerprint is SHA256:P0+5r+bZ6Gpc0sL2A0+kkZYBNHmjfucgne/ZBHPMHQ.
Are you sure you want to continue connecting (yes/no/[fingerprint])? yes
Warning: Permanently added 'ec2-13-234-38-84.ap-south-1.compute.amazonaws.com,13.234.38.84' (ECDSA) to the list
of known hosts.
Welcome to Ubuntu 24.04 LTS (GNU/Linux 6.8.0-1012-aws x86_64)

 * Documentation:  https://help.ubuntu.com
 * Management:    https://landscape.canonical.com
 * Support:       https://ubuntu.com/pro

System information as of Thu Sep 26 09:46:24 UTC 2024

System load:  0.08      Processes:    118
Usage of /:   22.8% of 6.71GB  Users logged in:  0
Memory usage: 5%      IPv4 address for enX0: 172.31.47.161
Swap usage:   0%

```

```

PS C:\Windows\system32> cd C:\Users\Dell\Desktop\keypair
PS C:\Users\Dell\Desktop\keypair> ssh -i "exp3key.pem" ubuntu@ec2-65-0-74-51.ap-south-1.compute.amazonaws.com
The authenticity of host 'ec2-65-0-74-51.ap-south-1.compute.amazonaws.com (65.0.74.51)' can't be established.
ECDSA key fingerprint is SHA256:RImw0JI2RUM5s3V2bYLL681AJE4iEasupz4bVTCjFw.
Are you sure you want to continue connecting (yes/no/[fingerprint])? yes
Warning: Permanently added 'ec2-65-0-74-51.ap-south-1.compute.amazonaws.com,65.0.74.51' (ECDSA) to the list of
known hosts.
Welcome to Ubuntu 24.04 LTS (GNU/Linux 6.8.0-1012-aws x86_64)

 * Documentation:  https://help.ubuntu.com
 * Management:    https://landscape.canonical.com
 * Support:       https://ubuntu.com/pro

System information as of Thu Sep 26 09:47:05 UTC 2024

System load:  0.13      Processes:    119
Usage of /:   22.8% of 6.71GB  Users logged in:  0
Memory usage: 5%      IPv4 address for enX0: 172.31.33.106
Swap usage:   0%

Expanded Security Maintenance for Applications is not enabled.

```

Step 3: Docker installation:

Run the following command: These commands are used to install Docker on an Ubuntu system by adding Docker's official GPG key and configuring the Docker repository.

```

curl -fsSL https://download.docker.com/linux/ubuntu/gpg | sudo apt-key add -
curl -fsSL https://download.docker.com/linux/ubuntu/gpg | sudo tee
/etc/apt/trusted.gpg.d/docker.gpg > /dev/null
sudo add-apt-repository "deb [arch=amd64] https://download.docker.com/linux/ubuntu
$(lsb_release -cs) stable"

```

```

ubuntu@ip-172-31-37-198:~$ curl -fsSL https://download.docker.com/linux/ubuntu/gpg | sudo tee /etc/apt/trusted.gpg.d/docker.gpg > /dev/null
ubuntu@ip-172-31-37-198:~$ sudo add-apt-repository "deb [arch=amd64] https://download.docker.com/linux/ubuntu $(lsb_release -cs) stable"
Repository: 'deb [arch=amd64] https://download.docker.com/linux/ubuntu noble stable'
Description:
Archive for codename: noble components: stable
More info: https://download.docker.com/linux/ubuntu
Adding repository.
Press [ENTER] to continue or Ctrl-C to cancel.
Adding deb entry to /etc/apt/sources.list.d/archive_uri-https_download_docker_com_linux_ubuntu-noble.list
Adding disabled deb-src entry to /etc/apt/sources.list.d/archive_uri-https_download_docker_com_linux_ubuntu-noble.list
Hit:1 http://ap-south-1.ec2.archive.ubuntu.com/ubuntu noble InRelease
Get:2 http://ap-south-1.ec2.archive.ubuntu.com/ubuntu noble-updates InRelease [126 kB]
Get:3 http://ap-south-1.ec2.archive.ubuntu.com/ubuntu noble-backports InRelease [126 kB]
Get:4 https://download.docker.com/linux/ubuntu noble InRelease [48.8 kB]
Get:5 http://ap-south-1.ec2.archive.ubuntu.com/ubuntu noble/universe amd64 Packages [15.0 MB]
Get:6 https://security.ubuntu.com/ubuntu noble-security InRelease [126 kB]
Get:7 http://ap-south-1.ec2.archive.ubuntu.com/ubuntu noble/universe Translation-en [5982 kB]
Get:8 https://download.docker.com/linux/ubuntu noble/stable amd64 Packages [15.3 kB]
Get:9 http://ap-south-1.ec2.archive.ubuntu.com/ubuntu noble/universe amd64 Components [3871 kB]
Get:10 http://ap-south-1.ec2.archive.ubuntu.com/ubuntu noble/universe amd64 c-n-f Metadata [301 kB]
Get:11 http://ap-south-1.ec2.archive.ubuntu.com/ubuntu noble/multiverse amd64 Packages [269 kB]
Get:12 http://ap-south-1.ec2.archive.ubuntu.com/ubuntu noble/multiverse Translation-en [118 kB]
Get:13 http://ap-south-1.ec2.archive.ubuntu.com/ubuntu noble/multiverse amd64 Components [35.0 kB]
Get:14 http://ap-south-1.ec2.archive.ubuntu.com/ubuntu noble/multiverse amd64 c-n-f Metadata [8328 B]

```

Run the following commands to refresh your local package list to ensure the latest packages are available and install Docker Community Edition on your system without prompting for confirmation.

```
sudo apt-get update
```

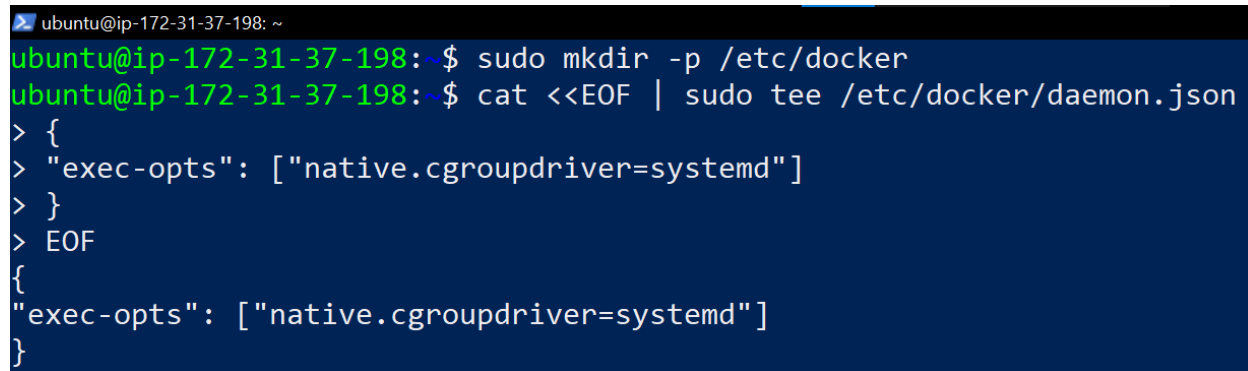
```
sudo apt-get install -y docker-ce
```

```
ubuntu@ip-172-31-37-198: ~  
ubuntu@ip-172-31-37-198: $ sudo apt-get update  
Hit:1 http://ap-south-1.ec2.archive.ubuntu.com/ubuntu noble InRelease  
Hit:2 http://ap-south-1.ec2.archive.ubuntu.com/ubuntu noble-updates InRelease  
Hit:3 http://ap-south-1.ec2.archive.ubuntu.com/ubuntu noble-backports InRelease  
Hit:4 https://download.docker.com/linux/ubuntu noble InRelease  
Hit:5 http://security.ubuntu.com/ubuntu noble-security InRelease  
Reading package lists... Done  
W: https://download.docker.com/linux/ubuntu/dists/noble/InRelease: The key(s) in the keyring /etc/apt/trusted.gpg.d/docker.gpg are ignored as the file has an unsupported filetype.  
W: https://download.docker.com/linux/ubuntu/dists/noble/InRelease: Key is stored in legacy trusted.gpg keyring (/etc/apt/trusted.gpg), see the DEPRECATION section in apt-key(8) for details.  
ubuntu@ip-172-31-37-198: $ sudo apt-get install -y docker-ce  
Reading package lists... Done  
Building dependency tree... Done  
Reading state information... Done  
The following additional packages will be installed:  
  containerd.io docker-buildx-plugin docker-ce-cli docker-ce-rootless-extras docker-compose-plugin libltdl7  
  libsllp0 pigz sllrp4netns  
Suggested packages:  
  aufs-tools cgroupfs-mount | cgroup-lite  
The following NEW packages will be installed:  
  containerd.io docker-buildx-plugin docker-ce docker-ce-cli docker-ce-rootless-extras docker-compose-plugin  
  libltdl7 libsllp0 pigz sllrp4netns  
0 upgraded, 10 newly installed, 0 to remove and 142 not upgraded.  
Need to get 123 MB of archives.  
After this operation, 442 MB of additional disk space will be used.  
Get:1 http://ap-south-1.ec2.archive.ubuntu.com/ubuntu noble/universe amd64 pigz amd64 2.8-1 [65.6 kB]
```

```
ubuntu@ip-172-31-37-198: ~  
Created symlink /etc/systemd/system/multi-user.target.wants/containerd.service → /usr/lib/systemd/system/containerd.service.  
Setting up docker-compose-plugin (2.29.7-1~ubuntu.24.04~noble) ...  
Setting up libltdl7:amd64 (2.4.7-7build1) ...  
Setting up docker-ce-cli (5:27.3.1-1~ubuntu.24.04~noble) ...  
Setting up libsllp0:amd64 (4.7.0-1ubuntu3) ...  
Setting up pigz (2.8-1) ...  
Setting up docker-ce-rootless-extras (5:27.3.1-1~ubuntu.24.04~noble) ...  
Setting up sllrp4netns (1.2.1-1build2) ...  
Setting up docker-ce (5:27.3.1-1~ubuntu.24.04~noble) ...  
Created symlink /etc/systemd/system/multi-user.target.wants/docker.service → /usr/lib/systemd/system/docker.service.  
Created symlink /etc/systemd/system/sockets.target.wants/docker.socket → /usr/lib/systemd/system/docker.socket.  
Processing triggers for man-db (2.12.0-4build2) ...  
Processing triggers for libc-bin (2.39-0ubuntu8.2) ...  
Scanning processes...  
Scanning linux images...  
  
Running kernel seems to be up-to-date.  
  
No services need to be restarted.  
  
No containers need to be restarted.  
  
No user sessions are running outdated binaries.  
  
No VM guests are running outdated hypervisor (qemu) binaries on this host.
```

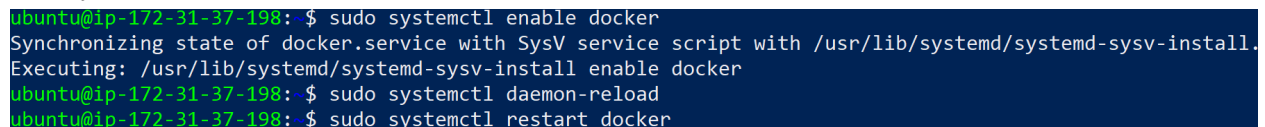
Run the following commands to create the Docker configuration directory and write a configuration file for Docker to use the systemd cgroup driver.

```
sudo mkdir -p /etc/docker
cat <<EOF | sudo tee /etc/docker/daemon.json
{
"exec-opts": ["native.cgroupdriver=systemd"]
}
EOF
```

A terminal window with a dark blue background and white text. The prompt is 'ubuntu@ip-172-31-37-198: ~'. The user enters 'sudo mkdir -p /etc/docker'. The prompt changes to 'ubuntu@ip-172-31-37-198:~\$'. The user enters 'cat <<EOF | sudo tee /etc/docker/daemon.json'. The terminal shows the input of a JSON configuration file: '{', '"exec-opts": ["native.cgroupdriver=systemd"]', '}', and 'EOF'. The prompt returns to 'ubuntu@ip-172-31-37-198:~\$'.

Run the following commands to ensure Docker starts automatically on system boot, reload systemd to apply new configurations, restart Docker to apply the new configuration.

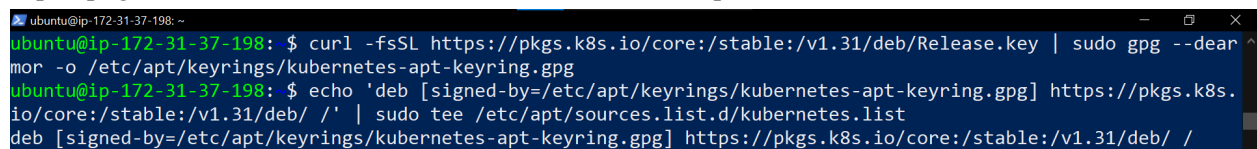
```
sudo systemctl enable docker
sudo systemctl daemon-reload
sudo systemctl restart docker
```

A terminal window with a dark blue background and white text. The prompt is 'ubuntu@ip-172-31-37-198: ~'. The user enters 'sudo systemctl enable docker'. The terminal shows the output: 'Synchronizing state of docker.service with SysV service script with /usr/lib/systemd/systemd-sysv-install. Executing: /usr/lib/systemd/systemd-sysv-install enable docker'. The prompt changes to 'ubuntu@ip-172-31-37-198:~\$'. The user enters 'sudo systemctl daemon-reload'. The prompt changes to 'ubuntu@ip-172-31-37-198:~\$'. The user enters 'sudo systemctl restart docker'.

Step 4: Kubernetes Installation:

Run the following commands to add the Kubernetes package repository to your Ubuntu system in order to install Kubernetes components like kubectl, kubelet, and kubeadm.

```
curl -fsSL https://pkgs.k8s.io/core:/stable:/v1.31/deb/Release.key | sudo gpg --dearmor -o
/etc/apt/keyrings/kubernetes-apt-keyring.gpg
echo 'deb [signed-by=/etc/apt/keyrings/kubernetes-apt-keyring.gpg]
https://pkgs.k8s.io/core:/stable:/v1.31/deb/ /' | sudo tee /etc/apt/sources.list.d/kubernetes.list
```

A terminal window with a dark blue background and white text. The prompt is 'ubuntu@ip-172-31-37-198: ~'. The user enters 'curl -fsSL https://pkgs.k8s.io/core:/stable:/v1.31/deb/Release.key | sudo gpg --dearmor -o /etc/apt/keyrings/kubernetes-apt-keyring.gpg'. The prompt changes to 'ubuntu@ip-172-31-37-198:~\$'. The user enters 'echo 'deb [signed-by=/etc/apt/keyrings/kubernetes-apt-keyring.gpg] https://pkgs.k8s.io/core:/stable:/v1.31/deb/ /' | sudo tee /etc/apt/sources.list.d/kubernetes.list'. The terminal shows the output of the echo command. The prompt returns to 'ubuntu@ip-172-31-37-198:~\$'.

Next, run these commands to refresh the local package index to include the newly added Kubernetes repository, install the main Kubernetes components (kubelet, kubeadm, kubectl), prevent the installed Kubernetes components from being automatically updated, ensuring cluster stability until manual updates are performed.

```
sudo apt-get update
```

```
sudo apt-get install -y kubelet kubeadm kubectl
```

```
sudo apt-mark hold kubelet kubeadm kubectl
```

```
ubuntu@ip-172-31-37-198: ~$ sudo apt-get update
Hit:1 http://ap-south-1.ec2.archive.ubuntu.com/ubuntu noble InRelease
Get:2 http://ap-south-1.ec2.archive.ubuntu.com/ubuntu noble-updates InRelease [126 kB]
Hit:3 http://ap-south-1.ec2.archive.ubuntu.com/ubuntu noble-backports InRelease
Hit:4 https://download.docker.com/linux/ubuntu noble InRelease
Get:5 http://ap-south-1.ec2.archive.ubuntu.com/ubuntu noble-updates/main amd64 Packages [533 kB]
Get:6 http://ap-south-1.ec2.archive.ubuntu.com/ubuntu noble-updates/main Translation-en [129 kB]
Get:7 http://ap-south-1.ec2.archive.ubuntu.com/ubuntu noble-updates/universe amd64 Packages [376 kB]
Hit:8 http://security.ubuntu.com/ubuntu noble-security InRelease
Get:9 http://ap-south-1.ec2.archive.ubuntu.com/ubuntu noble-updates/universe Translation-en [155 kB]
Get:10 https://prod-cdn.packages.k8s.io/repositories/isv/kubernetes:/core:/stable:/v1.31/deb InRelease [1186 B]
Get:11 https://prod-cdn.packages.k8s.io/repositories/isv/kubernetes:/core:/stable:/v1.31/deb Packages [4865 B]
Fetched 1324 kB in 1s (2118 kB/s)
Reading package lists... Done
W: https://download.docker.com/linux/ubuntu/dists/noble/InRelease: The key(s) in the keyring /etc/apt/trusted.gpg.d/docker.gpg are ignored as the file has an unsupported filetype.
W: https://download.docker.com/linux/ubuntu/dists/noble/InRelease: Key is stored in legacy trusted.gpg keyring (/etc/apt/trusted.gpg), see the DEPRECATION section in apt-key(8) for details.
ubuntu@ip-172-31-37-198: ~$ sudo apt-get install -y kubelet kubeadm kubectl
Reading package lists... Done
Building dependency tree... Done
Reading state information... Done
The following additional packages will be installed:
  conntack cri-tools kubernetes-cni
The following NEW packages will be installed:
```

```
Preparing to unpack .../3-kubectl_1.31.1-1.1_amd64.deb ...
Unpacking kubectl (1.31.1-1.1) ...
Selecting previously unselected package kubernetes-cni.
Preparing to unpack .../4-kubernetes-cni_1.5.1-1.1_amd64.deb ...
Unpacking kubernetes-cni (1.5.1-1.1) ...
Selecting previously unselected package kubelet.
Preparing to unpack .../5-kubelet_1.31.1-1.1_amd64.deb ...
Unpacking kubelet (1.31.1-1.1) ...
Setting up conntack (1:1.4.8-1ubuntu1) ...
Setting up kubectl (1.31.1-1.1) ...
Setting up cri-tools (1.31.1-1.1) ...
Setting up kubernetes-cni (1.5.1-1.1) ...
Setting up kubeadm (1.31.1-1.1) ...
Setting up kubelet (1.31.1-1.1) ...
Processing triggers for man-db (2.12.0-4build2) ...
Scanning processes...
Scanning linux images...

Running kernel seems to be up-to-date.

No services need to be restarted.

No containers need to be restarted.

No user sessions are running outdated binaries.

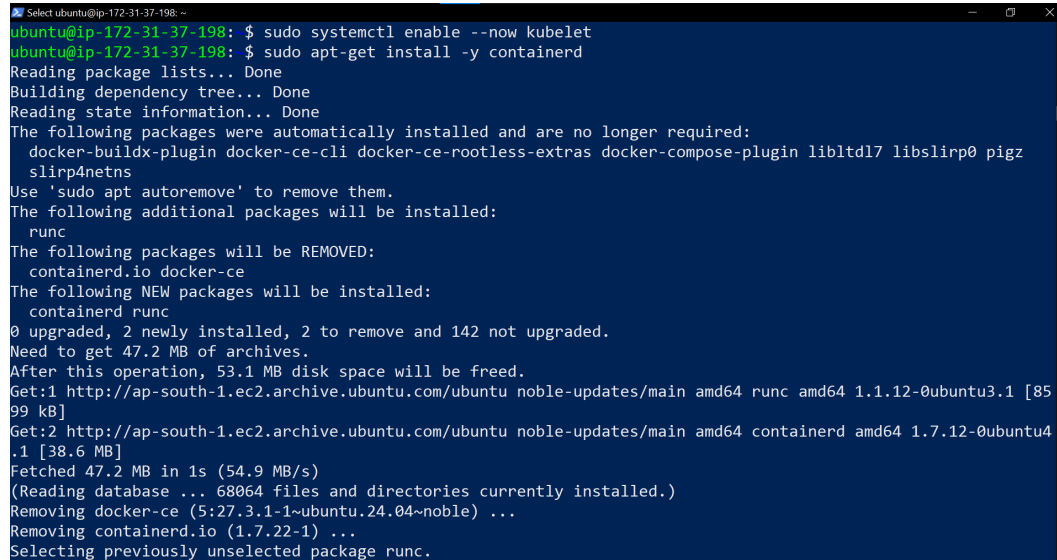
No VM guests are running outdated hypervisor (qemu) binaries on this host.
```

```
ubuntu@ip-172-31-37-198: ~$ sudo apt-mark hold kubelet kubeadm kubectl
kubelet set on hold.
kubeadm set on hold.
kubectl set on hold.
```


The command **sudo systemctl enable --now kubelet** enables the **kubelet** service, ensuring it starts automatically at boot and immediately starts running without needing a system reboot. The **sudo apt-get install -y containerd** command installs **containerd**, a container runtime that Kubernetes uses to manage and run containers.

```
sudo systemctl enable --now kubelet
```

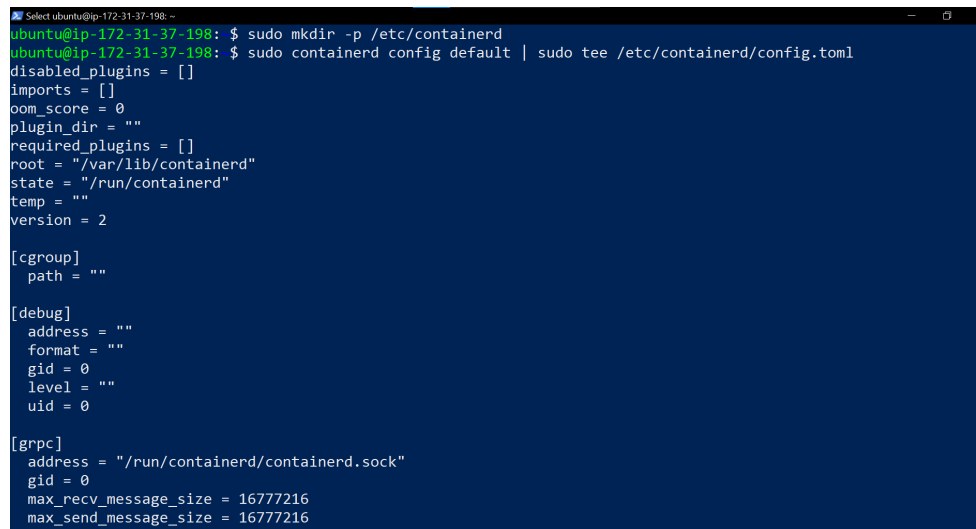
```
sudo apt-get install -y containerd
```

A terminal window showing the execution of two commands. The first command, 'sudo systemctl enable --now kubelet', is followed by the second command, 'sudo apt-get install -y containerd'. The output of the second command shows the package lists being read, the dependency tree being built, and state information being read. It then lists packages that were automatically installed and are no longer required, including docker-buildx-plugin, docker-ce-cli, docker-ce-rootless-extras, docker-compose-plugin, libltdl7, libslirp0, pigz, and slirp4netns. It suggests using 'sudo apt autoremove' to remove them. Next, it lists additional packages to be installed: runc. It then lists packages to be removed: containerd.io and docker-ce. Finally, it lists the new packages to be installed: containerd and runc. The output also shows that 0 packages are upgraded, 2 are newly installed, 2 are to be removed, and 142 are not upgraded. It indicates that 47.2 MB of archives are needed and that 53.1 MB of disk space will be freed after the operation. The output concludes with the fetching of archives and the removal of the old packages, ending with 'Selecting previously unselected package runc.'

The command `sudo mkdir -p /etc/containerd` creates the directory for containerd configuration files if it doesn't already exist. The next command, `sudo containerd config default | sudo tee /etc/containerd/config.toml`, generates a default configuration for containerd and saves it as `config.toml` in that directory. Together, these commands set up the necessary directory and create a default configuration file for containerd.

```
sudo mkdir -p /etc/containerd
```

```
sudo containerd config default | sudo tee /etc/containerd/config.toml
```

A terminal window showing the execution of two commands. The first command, 'sudo mkdir -p /etc/containerd', creates the directory. The second command, 'sudo containerd config default | sudo tee /etc/containerd/config.toml', generates the default configuration file. The output of the second command shows the configuration details for containerd, including disabled_plugins, imports, oom_score, plugin_dir, required_plugins, root, state, temp, and version. It also shows the configuration for the [cgroup], [debug], and [grpc] sections.

The command `sudo systemctl restart containerd` restarts the `containerd` service to apply any changes. `sudo systemctl enable containerd` sets it to start automatically at boot, while `sudo systemctl status containerd` checks its current status, showing whether it's active and any errors. Together, these commands manage the operation and health of the container runtime.

sudo systemctl restart containerd

sudo systemctl enable containerd

sudo systemctl status containerd

```
ubuntu@ip-172-31-37-198: $ sudo systemctl restart containerd
ubuntu@ip-172-31-37-198: $ sudo systemctl enable containerd
ubuntu@ip-172-31-37-198: $ sudo systemctl status containerd
● containerd.service - containerd container runtime
   Loaded: loaded (/usr/lib/systemd/system/containerd.service; enabled; preset: enabled)
   Active: active (running) since Thu 2024-09-26 10:09:37 UTC; 14s ago
     Docs: https://containerd.io
   Main PID: 4771 (containerd)
    Tasks: 7
   Memory: 13.0M (peak: 13.4M)
      CPU: 123ms
   CGroup: /system.slice/containerd.service
           └─4771 /usr/bin/containerd

Sep 26 10:09:37 ip-172-31-37-198 containerd[4771]: time="2024-09-26T10:09:37.127045263Z" level=info msg="Start>
Sep 26 10:09:37 ip-172-31-37-198 containerd[4771]: time="2024-09-26T10:09:37.127119958Z" level=info msg="Start>
Sep 26 10:09:37 ip-172-31-37-198 containerd[4771]: time="2024-09-26T10:09:37.127202179Z" level=info msg="Start>
Sep 26 10:09:37 ip-172-31-37-198 containerd[4771]: time="2024-09-26T10:09:37.127214694Z" level=info msg="Start>
Sep 26 10:09:37 ip-172-31-37-198 containerd[4771]: time="2024-09-26T10:09:37.127239102Z" level=info msg="Start>
Sep 26 10:09:37 ip-172-31-37-198 containerd[4771]: time="2024-09-26T10:09:37.127260747Z" level=info msg="Start>
Sep 26 10:09:37 ip-172-31-37-198 containerd[4771]: time="2024-09-26T10:09:37.127807399Z" level=info msg="servin>
Sep 26 10:09:37 ip-172-31-37-198 containerd[4771]: time="2024-09-26T10:09:37.127910878Z" level=info msg="servin>
Sep 26 10:09:37 ip-172-31-37-198 containerd[4771]: time="2024-09-26T10:09:37.128176100Z" level=info msg="contat>
Sep 26 10:09:37 ip-172-31-37-198 systemd[1]: Started containerd.service - containerd container runtime.
```

Socat installation:

sudo apt-get install -y socat

```
ubuntu@ip-172-31-37-198: $ sudo apt-get install -y socat
Reading package lists... Done
Building dependency tree... Done
Reading state information... Done
The following packages were automatically installed and are no longer required:
  docker-buildx-plugin docker-ce-cli docker-ce-rootless-extras docker-compose-plugin libltdl7 libslirp0 pigz
  slirp4netns
Use 'sudo apt autoremove' to remove them.
The following NEW packages will be installed:
  socat
0 upgraded, 1 newly installed, 0 to remove and 142 not upgraded.
Need to get 374 kB of archives.
After this operation, 1649 kB of additional disk space will be used.
Get:1 http://ap-south-1.ec2.archive.ubuntu.com/ubuntu noble/main amd64 socat amd64 1.8.0.0-4build3 [374 kB]
Fetched 374 kB in 0s (16.7 MB/s)
Selecting previously unselected package socat.
(Reading database ... 68108 files and directories currently installed.)
Preparing to unpack .../socat_1.8.0.0-4build3_amd64.deb ...
Unpacking socat (1.8.0.0-4build3) ...
Setting up socat (1.8.0.0-4build3) ...
Processing triggers for man-db (2.12.0-4build2) ...
Scanning processes...
Scanning linux images...

Running kernel seems to be up-to-date.

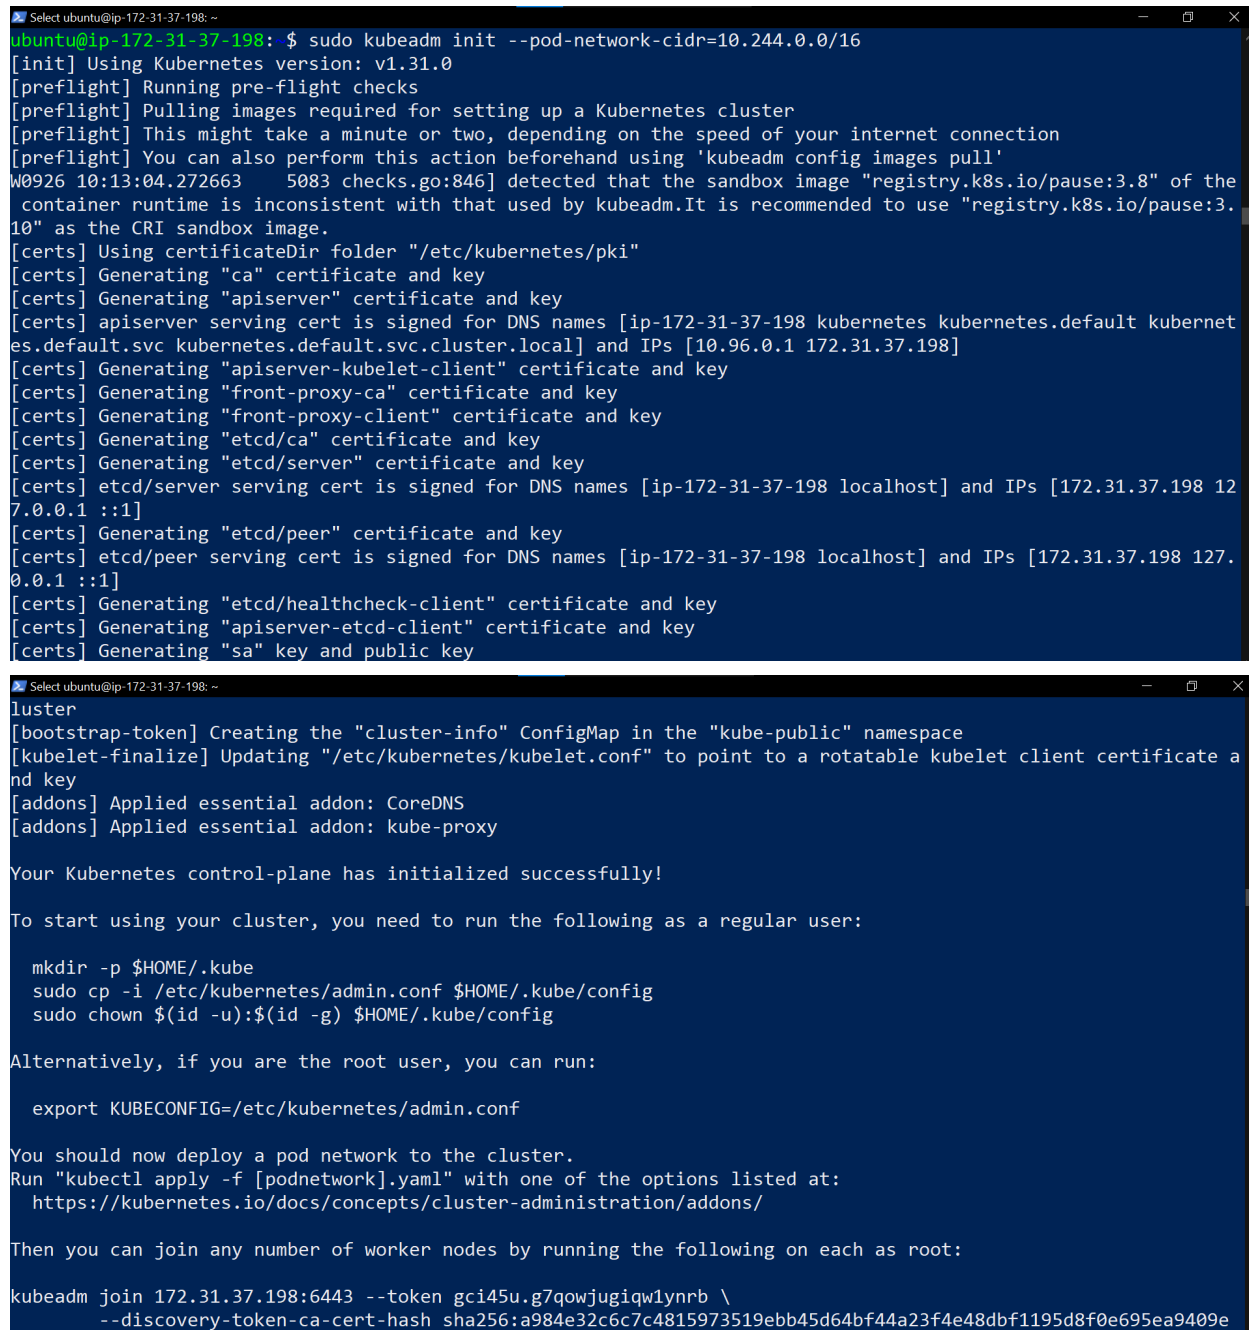
No services need to be restarted.
```

Step 5: Kubernetes Cluster

Run this command only on the master instance

```
sudo kubeadm init --pod-network-cidr=10.244.0.0/16
```

It initializes a Kubernetes cluster with kubeadm and sets up the control plane (master node).



```
Select ubuntu@ip-172-31-37-198: ~
ubuntu@ip-172-31-37-198: $ sudo kubeadm init --pod-network-cidr=10.244.0.0/16
[init] Using Kubernetes version: v1.31.0
[preflight] Running pre-flight checks
[preflight] Pulling images required for setting up a Kubernetes cluster
[preflight] This might take a minute or two, depending on the speed of your internet connection
[preflight] You can also perform this action beforehand using 'kubeadm config images pull'
W0926 10:13:04.272663    5083 checks.go:846] detected that the sandbox image "registry.k8s.io/pause:3.8" of the
container runtime is inconsistent with that used by kubeadm. It is recommended to use "registry.k8s.io/pause:3.
10" as the CRI sandbox image.
[certs] Using certificateDir folder "/etc/kubernetes/pki"
[certs] Generating "ca" certificate and key
[certs] Generating "apiserver" certificate and key
[certs] apiserver serving cert is signed for DNS names [ip-172-31-37-198 kubernet es.default kubernet
es.default.svc kubernet es.default.svc.cluster.local] and IPs [10.96.0.1 172.31.37.198]
[certs] Generating "apiserver-kubelet-client" certificate and key
[certs] Generating "front-proxy-ca" certificate and key
[certs] Generating "front-proxy-client" certificate and key
[certs] Generating "etcd/ca" certificate and key
[certs] Generating "etcd/server" certificate and key
[certs] etcd/server serving cert is signed for DNS names [ip-172-31-37-198 localhost] and IPs [172.31.37.198 12
7.0.0.1 ::1]
[certs] Generating "etcd/peer" certificate and key
[certs] etcd/peer serving cert is signed for DNS names [ip-172-31-37-198 localhost] and IPs [172.31.37.198 127.
0.0.1 ::1]
[certs] Generating "etcd/healthcheck-client" certificate and key
[certs] Generating "apiserver-etcd-client" certificate and key
[certs] Generating "sa" key and public key

cluster
[bootstrap-token] Creating the "cluster-info" ConfigMap in the "kube-public" namespace
[kubelet-finalize] Updating "/etc/kubernetes/kubelet.conf" to point to a rotatable kubelet client certificate a
nd key
[addons] Applied essential addon: CoreDNS
[addons] Applied essential addon: kube-proxy

Your Kubernetes control-plane has initialized successfully!

To start using your cluster, you need to run the following as a regular user:

  mkdir -p $HOME/.kube
  sudo cp -i /etc/kubernetes/admin.conf $HOME/.kube/config
  sudo chown $(id -u):$(id -g) $HOME/.kube/config

Alternatively, if you are the root user, you can run:

  export KUBECONFIG=/etc/kubernetes/admin.conf

You should now deploy a pod network to the cluster.
Run "kubectl apply -f [podnetwork].yaml" with one of the options listed at:
  https://kubernetes.io/docs/concepts/cluster-administration/addons/

Then you can join any number of worker nodes by running the following on each as root:

kubeadm join 172.31.37.198:6443 --token gci45u.g7qowjugiqw1ynrb \
  --discovery-token-ca-cert-hash sha256:a984e32c6c7c4815973519ebb45d64bf44a23f4e48dbf1195d8f0e695ea9409e
```

Run this command on master and also copy and save the Join command from above.

```
mkdir -p $HOME/.kube
```

```
sudo cp -i /etc/kubernetes/admin.conf $HOME/.kube/config
```

```
sudo chown $(id -u):$(id -g) $HOME/.kube/config
```

```
ubuntu@ip-172-31-37-198:~$ mkdir -p $HOME/.kube
ubuntu@ip-172-31-37-198:~$ sudo cp -i /etc/kubernetes/admin.conf $HOME/.kube/config
ubuntu@ip-172-31-37-198:~$ sudo chown $(id -u):$(id -g) $HOME/.kube/config
ubuntu@ip-172-31-37-198:~$ kubectl get nodes
```

NAME	STATUS	ROLES	AGE	VERSION
ip-172-31-37-198	NotReady	control-plane	12m	v1.31.1

Connect master and worker nodes by running this command on the worker logged in terminals:

I ran the following command to do this,

```
kubeadm join 172.31.37.198:6443 --token gci45u.g7qowjugiqwlynr \ --discovery-token-ca-cert-hash sha256:a984e32c6c7c4815973519ebb45d64bf44a23f4e48dbf1195d8f0e695ea9409e
```

On node 1:

```
ubuntu@ip-172-31-47-161:~$ sudo kubeadm join 172.31.37.198:6443 --token gci45u.g7qowjugiqwlynr --discovery-token-ca-cert-hash sha256:a984e32c6c7c4815973519ebb45d64bf44a23f4e48dbf1195d8f0e695ea9409e
[preflight] Running pre-flight checks
[preflight] Reading configuration from the cluster...
[preflight] FYI: You can look at this config file with 'kubectl -n kube-system get cm kubeadm-config -o yaml'
[kubelet-start] Writing kubelet configuration to file "/var/lib/kubelet/config.yaml"
[kubelet-start] Writing kubelet environment file with flags to file "/var/lib/kubelet/kubeadm-flags.env"
[kubelet-start] Starting the kubelet
[kubelet-check] Waiting for a healthy kubelet at http://127.0.0.1:10248/healthz. This can take up to 4m0s
[kubelet-check] The kubelet is healthy after 1.001901347s
[kubelet-start] Waiting for the kubelet to perform the TLS Bootstrap

This node has joined the cluster:
* Certificate signing request was sent to apiserver and a response was received.
* The Kubelet was informed of the new secure connection details.

Run 'kubectl get nodes' on the control-plane to see this node join the cluster.

ubuntu@ip-172-31-47-161:~$
```

On node 2:

```
ubuntu@ip-172-31-33-106:~$ sudo kubeadm join 172.31.37.198:6443 --token gci45u.g7qowjugiqwlynr --discovery-token-ca-cert-hash sha256:a984e32c6c7c4815973519ebb45d64bf44a23f4e48dbf1195d8f0e695ea9409e
[preflight] Running pre-flight checks
[preflight] Reading configuration from the cluster...
[preflight] FYI: You can look at this config file with 'kubectl -n kube-system get cm kubeadm-config -o yaml'
[kubelet-start] Writing kubelet configuration to file "/var/lib/kubelet/config.yaml"
[kubelet-start] Writing kubelet environment file with flags to file "/var/lib/kubelet/kubeadm-flags.env"
[kubelet-start] Starting the kubelet
[kubelet-check] Waiting for a healthy kubelet at http://127.0.0.1:10248/healthz. This can take up to 4m0s
[kubelet-check] The kubelet is healthy after 1.002345051s
[kubelet-start] Waiting for the kubelet to perform the TLS Bootstrap

This node has joined the cluster:
* Certificate signing request was sent to apiserver and a response was received.
* The Kubelet was informed of the new secure connection details.

Run 'kubectl get nodes' on the control-plane to see this node join the cluster.

ubuntu@ip-172-31-33-106:~$
```

Run kubectl get nodes on the master instance to confirm the joining of the worker nodes

```
Select ubuntu@ip-172-31-37-198: ~
ubuntu@ip-172-31-37-198:~$ kubectl get nodes
```

NAME	STATUS	ROLES	AGE	VERSION
ip-172-31-33-106	NotReady	<none>	17s	v1.31.1
ip-172-31-37-198	NotReady	control-plane	18m	v1.31.1
ip-172-31-47-161	NotReady	<none>	39s	v1.31.1

Since Status is NotReady we have to add a network plugin. And also we have to give the name to the nodes.

kubectl apply -f <https://docs.projectcalico.org/manifests/calico.yaml>

The command `kubectl apply -f https://docs.projectcalico.org/manifests/calico.yaml` deploys Calico, a networking and network security solution for Kubernetes, by applying the configuration specified in the provided YAML file from the Calico documentation. This sets up the necessary resources to enable networking capabilities within the Kubernetes cluster.

```
Select ubuntu@ip-172-31-37-198: ~
ubuntu@ip-172-31-37-198: ~$ kubectl apply -f https://docs.projectcalico.org/manifests/calico.yaml
poddisruptionbudget.policy/calico-kube-controllers created
serviceaccount/calico-kube-controllers created
serviceaccount/calico-node created
configmap/calico-config created
customresourcedefinition.apiextensions.k8s.io/bgpconfigurations.crd.projectcalico.org created
customresourcedefinition.apiextensions.k8s.io/bgppeers.crd.projectcalico.org created
customresourcedefinition.apiextensions.k8s.io/blockaffinities.crd.projectcalico.org created
customresourcedefinition.apiextensions.k8s.io/caliconodestatuses.crd.projectcalico.org created
customresourcedefinition.apiextensions.k8s.io/clusterinformations.crd.projectcalico.org created
customresourcedefinition.apiextensions.k8s.io/felixconfigurations.crd.projectcalico.org created
customresourcedefinition.apiextensions.k8s.io/globalnetworkpolicies.crd.projectcalico.org created
customresourcedefinition.apiextensions.k8s.io/globalnetworksets.crd.projectcalico.org created
customresourcedefinition.apiextensions.k8s.io/hostendpoints.crd.projectcalico.org created
customresourcedefinition.apiextensions.k8s.io/ipamblocks.crd.projectcalico.org created
customresourcedefinition.apiextensions.k8s.io/ipamconfigs.crd.projectcalico.org created
customresourcedefinition.apiextensions.k8s.io/ipamhandles.crd.projectcalico.org created
customresourcedefinition.apiextensions.k8s.io/ippools.crd.projectcalico.org created
customresourcedefinition.apiextensions.k8s.io/ipreservations.crd.projectcalico.org created
customresourcedefinition.apiextensions.k8s.io/kubecontrollersconfigurations.crd.projectcalico.org created
customresourcedefinition.apiextensions.k8s.io/networkpolicies.crd.projectcalico.org created
customresourcedefinition.apiextensions.k8s.io/networksets.crd.projectcalico.org created
clusterrole.rbac.authorization.k8s.io/calico-kube-controllers created
clusterrole.rbac.authorization.k8s.io/calico-node created
clusterrolebinding.rbac.authorization.k8s.io/calico-kube-controllers created
clusterrolebinding.rbac.authorization.k8s.io/calico-node created
daemonset.apps/calico-node created
```

`sudo systemctl status kubelet`

```
Select ubuntu@ip-172-31-37-198: ~
ubuntu@ip-172-31-37-198: ~$ sudo systemctl status kubelet
● kubelet.service - kubelet: The Kubernetes Node Agent
   Loaded: loaded (/usr/lib/systemd/system/kubelet.service; enabled; preset: enabled)
   Drop-In: /usr/lib/systemd/system/kubelet.service.d
            └─10-kubeadm.conf
   Active: active (running) since Thu 2024-09-26 10:13:42 UTC; 18min ago
     Docs: https://kubernetes.io/docs/
   Main PID: 5773 (kubelet)
    Tasks: 10 (limit: 4676)
   Memory: 32.6M (peak: 33.3M)
      CPU: 18.442s
   CGroup: /system.slice/kubelet.service
            └─5773 /usr/bin/kubelet --bootstrap-kubeconfig=/etc/kubernetes/bootstrap-kubelet.conf --kubeconfi

Sep 26 10:32:24 ip-172-31-37-198 kubelet[5773]: I0926 10:32:24.596915 5773 reconciler_common.go:245] "operat
Sep 26 10:32:24 ip-172-31-37-198 kubelet[5773]: I0926 10:32:24.596938 5773 reconciler_common.go:245] "operat
Sep 26 10:32:24 ip-172-31-37-198 kubelet[5773]: I0926 10:32:24.596958 5773 reconciler_common.go:245] "operat
Sep 26 10:32:24 ip-172-31-37-198 kubelet[5773]: I0926 10:32:24.596978 5773 reconciler_common.go:245] "operat
Sep 26 10:32:24 ip-172-31-37-198 kubelet[5773]: I0926 10:32:24.596997 5773 reconciler_common.go:245] "operat
Sep 26 10:32:24 ip-172-31-37-198 kubelet[5773]: I0926 10:32:24.597017 5773 reconciler_common.go:245] "operat
Sep 26 10:32:25 ip-172-31-37-198 kubelet[5773]: I0926 10:32:25.193047 5773 scope.go:117] "RemoveContainer" >
Sep 26 10:32:25 ip-172-31-37-198 kubelet[5773]: E0926 10:32:25.193206 5773 pod_workers.go:1301] "Error sync
Sep 26 10:32:27 ip-172-31-37-198 kubelet[5773]: E0926 10:32:27.544395 5773 kubelet.go:2902] "Container runti
Sep 26 10:32:32 ip-172-31-37-198 kubelet[5773]: E0926 10:32:32.545352 5773 kubelet.go:2902] "Container runti
lines 1-23/23 (END)
ubuntu@ip-172-31-37-198: ~$
```

`kubectl get nodes -o wide` helps us get to know that the Status is ready.

```

ubuntu@ip-172-31-37-198:~$ kubectl get nodes -o wide
NAME                                STATUS    ROLES    AGE   VERSION   INTERNAL-IP   EXTERNAL-IP   OS-IMAGE             KERNEL-VERSION   CONTAINER-RUNTIME
ip-172-31-33-106                    Ready     <none>    66s   v1.31.1   172.31.33.106 <none>        Ubuntu 24.04 LTS     6.8.0-1012-aws   containerd://1.7.12
ip-172-31-37-198                    Ready     control-plane 19m   v1.31.1   172.31.37.198 <none>        Ubuntu 24.04 LTS     6.8.0-1012-aws   containerd://1.7.12
ip-172-31-47-161                    Ready     <none>    88s   v1.31.1   172.31.47.161 <none>        Ubuntu 24.04 LTS     6.8.0-1012-aws   containerd://1.7.12
ubuntu@ip-172-31-37-198:~$ kubectl label node ip-172-31-28-117 kubernetes.io/role=Node1
Error from server (NotFound): nodes "ip-172-31-28-117" not found
ubuntu@ip-172-31-37-198:~$ kubectl label node ip-172-31-47-161 kubernetes.io/role=Node1
node/ip-172-31-47-161 labeled
ubuntu@ip-172-31-37-198:~$ kubectl label node ip-172-31-33-106 kubernetes.io/role=Node2
node/ip-172-31-33-106 labeled
ubuntu@ip-172-31-37-198:~$ kubectl get nodes -o wide
NAME                                STATUS    ROLES    AGE   VERSION   INTERNAL-IP   EXTERNAL-IP   OS-IMAGE             KERNEL-VERSION   CONTAINER-RUNTIME
ip-172-31-33-106                    Ready     Node2    3m26s v1.31.1   172.31.33.106 <none>        Ubuntu 24.04 LTS     6.8.0-1012-aws   containerd://1.7.12
ip-172-31-37-198                    Ready     control-plane 21m   v1.31.1   172.31.37.198 <none>        Ubuntu 24.04 LTS     6.8.0-1012-aws   containerd://1.7.12
ip-172-31-47-161                    Ready     Node1    3m48s v1.31.1   172.31.47.161 <none>        Ubuntu 24.04 LTS     6.8.0-1012-aws   containerd://1.7.12

```

Or else run kubectl get nodes command

```

ubuntu@ip-172-31-37-198:~$ kubectl get nodes
NAME                                STATUS    ROLES    AGE   VERSION
ip-172-31-33-106                    Ready     Node2    3m42s v1.31.1
ip-172-31-37-198                    Ready     control-plane 21m   v1.31.1
ip-172-31-47-161                    Ready     Node1    4m4s   v1.31.1
ubuntu@ip-172-31-37-198:~$

```

Conclusion: In this experiment, we successfully set up a Kubernetes cluster on AWS by creating the necessary infrastructure and configuring security groups. We installed Docker and Kubernetes components, initialized the master node, and connected the worker nodes. By deploying Calico for networking, we enabled effective communication within the cluster. This hands-on experience enhanced our understanding of Kubernetes architecture and equipped us with practical skills for managing containerized applications in a cloud environment.