

**Adv DevOps Lab Exp 03**

**Aim:** To understand the Kubernetes Cluster Architecture, install and Spin Up a Kubernetes Cluster on Linux Machines/Cloud

**Step 1:** Create 3 EC2 instances (1 master and 2 slaves). Select SSH option in the inbound rules. Created a key pair to be used commonly between all 3 instances created.

I selected AWS Linux as my operating system and enabled **t2 medium** option for kubernetes cluster to run smoothly

Instances (3) Info <span>Last updated 16 minutes ago</span> <span>Connect</span> <span>Instance state ▼</span> <span>Actions ▼</span> <span>Launch instances ▼</span>						
<input type="text" value="Find Instance by attribute or tag (case-sensitive)"/>				All states ▼		
<input type="checkbox"/>	Name <span>↗</span> ▼	Instance ID	Instance state ▼	Instance type ▼	Status check	Alarm status
<input type="checkbox"/>	Master	i-0589bc6d5e452907c	Running <span>🔍</span> <span>🔍</span>	t2.medium	2/2 checks passed	<a href="#">View alarms</a> +
<input type="checkbox"/>	Slave2	i-00856c9b137f44cf6	Running <span>🔍</span> <span>🔍</span>	t2.medium	2/2 checks passed	<a href="#">View alarms</a> +
<input type="checkbox"/>	Slave1	i-0beaf1d570ab06d73	Running <span>🔍</span> <span>🔍</span>	t2.medium	2/2 checks passed	<a href="#">View alarms</a> +

**Step 2:** Open git bash. Change your directory to Downloads and run chmod command on the key pair file that we created for all the EC2 instances that we launched in the earlier step. Hers, i have created a key pair with the name “ec2user” and gave it an extension .pem.

After assigning the key pair file the appropriate permissions, run the ssh command in the following format:

ssh -i <key pair.pem> ubuntu@<public ip address of the instance(DNS)>

Perform this command on all the 3 instances

```
Dell@DESKTOP-PSTUV9S MINGW64 ~ (main)
$ cd Downloads

Dell@DESKTOP-PSTUV9S MINGW64 ~/Downloads (main)
$ chmod 400 "ec2user.pem"
```

```
Dell@DESKTOP-PSTUV9S MINGW64 ~/Downloads (main)
$ ssh -i ec2user.pem ubuntu@ec2-35-175-113-217.compute-1.amazonaws.com
The authenticity of host 'ec2-35-175-113-217.compute-1.amazonaws.com (35.175.113.217)' can't be established.
ED25519 key fingerprint is SHA256:ahGNOA8a2dmhd/81hRe2a2C6tyvwt2t0bNhpR5/PJM.
This key is not known by any other names.
Are you sure you want to continue connecting (yes/no/[fingerprint])? yes
Warning: Permanently added 'ec2-35-175-113-217.compute-1.amazonaws.com' (ED25519) to the list of known hosts.
Welcome to Ubuntu 24.04 LTS (GNU/Linux 6.8.0-1012-aws x86_64)

 * Documentation:  https://help.ubuntu.com
 * Management:    https://landscape.canonical.com
 * Support:       https://ubuntu.com/pro

System information as of Sat Sep 14 11:31:16 UTC 2024

System load:  0.0          Processes:    104
Usage of /:   22.8% of 6.71GB  Users logged in: 0
Memory usage: 19%          IPv4 address for enx0: 172.31.71.82
Swap usage:   0%

Expanded Security Maintenance for Applications is not enabled.

0 updates can be applied immediately.

Enable ESM Apps to receive additional future security updates.
See https://ubuntu.com/esm or run: sudo pro status

The list of available updates is more than a week old.
To check for new updates run: sudo apt update

The programs included with the Ubuntu system are free software;
the exact distribution terms for each program are described in the
individual files in /usr/share/doc/*/copyright.

Ubuntu comes with ABSOLUTELY NO WARRANTY, to the extent permitted by
applicable law.

To run a command as administrator (user "root"), use "sudo <command>".
See "man sudo_root" for details.
```

**Step 3: Docker Installation**

Perform this step on all 3 instances

```
[ec2-user@ip-172-31-22-31 ~]$ yum install docker -y
Error: This command has to be run with superuser privileges (under the root user on most systems).
[ec2-user@ip-172-31-22-31 ~]$ sudo su
[root@ip-172-31-22-31 ec2-user]# yum install docker -y
Last metadata expiration check: 0:03:24 ago on Sat Sep 14 14:54:57 2024.
Dependencies resolved.
```

Package	Architecture	Version	Repository	Size
<b>Installing:</b>				
docker	x86_64	25.0.6-1.amzn2023.0.2	amazonlinux	44 M
<b>Installing dependencies:</b>				
containerd	x86_64	1.7.20-1.amzn2023.0.1	amazonlinux	35 M
iptables-libs	x86_64	1.8.8-3.amzn2023.0.2	amazonlinux	401 k
iptables-nft	x86_64	1.8.8-3.amzn2023.0.2	amazonlinux	183 k
libcgroup	x86_64	3.0-1.amzn2023.0.1	amazonlinux	75 k
libnetfilter_conntrack	x86_64	1.0.8-2.amzn2023.0.2	amazonlinux	58 k
libnftnl	x86_64	1.0.1-19.amzn2023.0.2	amazonlinux	30 k
libnftnl	x86_64	1.2.2-2.amzn2023.0.2	amazonlinux	84 k
pkgconf	x86_64	2.5-1.amzn2023.0.3	amazonlinux	83 k
runit	x86_64	1.1.13-1.amzn2023.0.1	amazonlinux	3.2 M

```
Transaction Summary
Install 10 Packages
Total download size: 84 M
Installed size: 317 M
```

Next, we are supposed to configure cgroup in a daemon.json file.

Run the following commands

```
cd /etc/docker
cat <<EOF | sudo tee /etc/docker/daemon.json
{
"exec-opts": ["native.cgroupdriver=systemd"],
"log-driver": "json-file",
"log-opts": {
"max-size": "100m"
},
"storage-driver": "overlay2"
}
EOF
```

sudo systemctl enable docker

sudo systemctl daemon-reload

sudo systemctl restart docker

```
[root@ip-172-31-21-0 ec2-user]# cd /etc/docker
[root@ip-172-31-21-0 docker]# cat <<EOF | sudo tee /etc/docker/daemon.json
{
"exec-opts": ["native.cgroupdriver=systemd"]
}
EOF
[root@ip-172-31-21-0 docker]# sudo systemctl enable docker
Created symlink /etc/systemd/system/multi-user.target.wants/docker.service → /usr/lib/systemd/system/docker.service.
[root@ip-172-31-21-0 docker]# sudo systemctl daemon-reload
[root@ip-172-31-21-0 docker]# sudo systemctl restart docker
[root@ip-172-31-21-0 docker]# # Set SELinux in permissive mode (effectively disabling it)
sudo setenforce 0
sudo sed -i 's/^SELINUX=enforcing$/SELINUX=permissive/' /etc/selinux/config
[root@ip-172-31-21-0 docker]# # This overwrites any existing configuration in /etc/yum.repos.d/kubernetes.repo
cat <<EOF | sudo tee /etc/yum.repos.d/kubernetes.repo
[kubernetes]
name=Kubernetes
baseurl=https://pkgs.k8s.io/core/stable/v1.31/rpm/
enabled=1
```

**Step 4: Kubernetes Installation**

Carry out this step on all 3 instances

Now, there are a number of steps to install kubernetes onto our instances

1. Set SELinux to permissive mode. These instructions are for Kubernetes 1.31.

# Set SELinux in permissive mode (effectively disabling it)

```
sudo setenforce 0
```

```
sudo sed -i 's/^SELINUX=enforcing$/SELINUX=permissive/' /etc/selinux/config
```

2. Add the Kubernetes yum repository. The exclude parameter in the repository definition ensures that the packages related to Kubernetes are not upgraded upon running yum update as there's a special procedure that must be followed for upgrading Kubernetes. Please note that this repository have packages only for Kubernetes 1.31; for other Kubernetes minor versions, you need to change the Kubernetes minor version in the URL to match your desired minor version (you should also check that you are reading the documentation for the version of Kubernetes that you plan to install).

# This overwrites any existing configuration in /etc/yum.repos.d/kubernetes.repo

```
cat <<EOF | sudo tee /etc/yum.repos.d/kubernetes.repo
```

```
[kubernetes]
```

```
name=Kubernetes
```

```
baseurl=https://pkgs.k8s.io/core:/stable:/v1.31/rpm/
```

```
enabled=1
```

```
gpgcheck=1
```

```
gpgkey=https://pkgs.k8s.io/core:/stable:/v1.31/rpm/repodata/repomd.xml.key
```

```
exclude=kubelet kubeadm kubectl cri-tools kubernetes-cni
```

```
EOF
```

3. Install kubelet, kubeadm and kubectl:

```
sudo yum install -y kubelet kubeadm kubectl --disableexcludes=kubernetes
```

4. (Optional) Enable the kubelet service before running kubeadm:

```
sudo systemctl enable --now kubelet
```

```
[root@ip-172-31-22-31 ec2-user]# # Set SELinux in permissive mode (effectively disabling it)
sudo setenforce 0
sudo sed -i 's/^SELINUX=enforcing$/SELINUX=permissive/' /etc/selinux/config
[root@ip-172-31-22-31 ec2-user]# # This overwrites any existing configuration in /etc/yum.repos.d/kubernetes.repo
cat <<EOF | sudo tee /etc/yum.repos.d/kubernetes.repo
[kubernetes]
name=Kubernetes
baseurl=https://pkgs.k8s.io/core:/stable:/v1.31/rpm/
enabled=1
gpgcheck=1
gpgkey=https://pkgs.k8s.io/core:/stable:/v1.31/rpm/repodata/repomd.xml.key
exclude=kubelet kubeadm kubectl cri-tools kubernetes-cni
EOF
[kubernetes]
name=Kubernetes
baseurl=https://pkgs.k8s.io/core:/stable:/v1.31/rpm/
enabled=1
gpgcheck=1
gpgkey=https://pkgs.k8s.io/core:/stable:/v1.31/rpm/repodata/repomd.xml.key
exclude=kubelet kubeadm kubectl cri-tools kubernetes-cni
[root@ip-172-31-22-31 ec2-user]# sudo yum install -y kubelet kubeadm kubectl --disableexcludes=kubernetes
Kubernetes                                         48 kB/s | 9.4 kB    00:00
Dependencies resolved.
=====
Package                               Architecture      Version           Repository      Size
=====
Installing:
kubeadm                               x86_64            1.31.1-150500.1.1 kubernetes     11 M
```

```

Installing      : conntrack-tools-1.4.6-2.amzn2023.0.2.x86_64      6/9
Running scriptlet: conntrack-tools-1.4.6-2.amzn2023.0.2.x86_64      6/9
Installing      : kubelet-1.31.1-150500.1.1.x86_64                7/9
Running scriptlet: kubelet-1.31.1-150500.1.1.x86_64                7/9
Installing      : kubeadm-1.31.1-150500.1.1.x86_64                8/9
Installing      : kubect1-1.31.1-150500.1.1.x86_64                9/9
Running scriptlet: kubect1-1.31.1-150500.1.1.x86_64                9/9
Verifying       : conntrack-tools-1.4.6-2.amzn2023.0.2.x86_64      1/9
Verifying       : libnetfilter_cthelper-1.0.0-21.amzn2023.0.2.x86_64 2/9
Verifying       : libnetfilter_cttimeout-1.0.0-19.amzn2023.0.2.x86_64 3/9
Verifying       : libnetfilter_queue-1.0.5-2.amzn2023.0.2.x86_64 4/9
Verifying       : cri-tools-1.31.1-150500.1.1.x86_64             5/9
Verifying       : kubeadm-1.31.1-150500.1.1.x86_64             6/9
Verifying       : kubect1-1.31.1-150500.1.1.x86_64             7/9
Verifying       : kubelet-1.31.1-150500.1.1.x86_64             8/9
Verifying       : kubernetes-cni-1.5.1-150500.1.1.x86_64        9/9

Installed:
  conntrack-tools-1.4.6-2.amzn2023.0.2.x86_64      cri-tools-1.31.1-150500.1.1.x86_64
  kubeadm-1.31.1-150500.1.1.x86_64                kubect1-1.31.1-150500.1.1.x86_64
  kubelet-1.31.1-150500.1.1.x86_64                kubernetes-cni-1.5.1-150500.1.1.x86_64
  libnetfilter_cthelper-1.0.0-21.amzn2023.0.2.x86_64  libnetfilter_cttimeout-1.0.0-19.amzn2023.0.2.x86_64
  libnetfilter_queue-1.0.5-2.amzn2023.0.2.x86_64

Complete!
[root@ip-172-31-22-31 ec2-user]# sudo systemctl enable --now kubelet
Created symlink /etc/systemd/system/multi-user.target.wants/kubelet.service → /usr/lib/systemd/system/kubelet.service.
[root@ip-172-31-22-31 ec2-user]#

```

After installing Kubernetes, we need to configure internet options to allow bridging.

- Sudo swapoff-a
- echo "net.bridge.bridge-nf-call-iptables=1"|sudo tee -a /etc/sysctl.conf
- Sudo sysctl-p

## Step 5:

### On master machine

Run command ...kubeadm init with the proper network pod, here it is, --pod-network-cidr=10.244.0.0/16 to initialize kubernetes

```

[bootstrap-token] Configured RBAC rules to allow the csrapprover controller automatically approve CSRs from a Node Bootstrap Token
[bootstrap-token] Configured RBAC rules to allow certificate rotation for all node client certificates in the cluster
[bootstrap-token] Creating the "cluster-info" ConfigMap in the "kube-public" namespace
[kubelet-finalize] Updating "/etc/kubernetes/kubelet.conf" to point to a rotatable kubelet client certificate and key
[addons] Applied essential addon: CoreDNS
[addons] Applied essential addon: kube-proxy

Your Kubernetes control-plane has initialized successfully!

To start using your cluster, you need to run the following as a regular user:

mkdir -p $HOME/.kube
sudo cp -i /etc/kubernetes/admin.conf $HOME/.kube/config
sudo chown $(id -u):$(id -g) $HOME/.kube/config

Alternatively, if you are the root user, you can run:

export KUBECONFIG=/etc/kubernetes/admin.conf

You should now deploy a pod network to the cluster.
Run "kubect1 apply -f [podnetwork].yaml" with one of the options listed at:
https://kubernetes.io/docs/concepts/cluster-administration/addons/

Then you can join any number of worker nodes by running the following on each as root:

kubeadm join 172.31.22.31:6443 --token plmkzy.2ocxer4410uwlqk4 \
--discovery-token-ca-cert-hash sha256:2590fd7ba571e7e92b4f18f77c2149583f19f6049e3dfb4d306ac22cf2f465d6
[root@ip-172-31-22-31 ec2-user]#

```

We are supposed to add a networking plugin named flannel with the help of the command mentioned in the console output

i.e

kubect1 apply -f <https://raw.githubusercontent.com/coreos/flannel/master/Documentation/kube-flannel.yml>

```

[root@ip-172-31-29-225 ec2-user]# kubect1 apply -f https://raw.githubusercontent.com/coreos/flannel/master/Documentation/kube-flannel
.yml
namespace/kube-flannel created
clusterrole.rbac.authorization.k8s.io/flannel created
clusterrolebinding.rbac.authorization.k8s.io/flannel created
serviceaccount/flannel created
configmap/kube-flannel-cfg created
daemonset.apps/kube-flannel-ds created

```

**On worker machines**

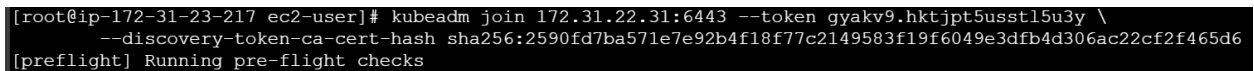
Run the following commands to ensure a smooth and secure joining to the master node

Paste the below command on all 2 worker machines

- `sudo yum install iproute-tc-y`
- `sudo systemctl enable kubelet`
- `sudo systemctl restart kubelet`

Then we are supposed to run the join command that was generated in the console output of our master machine

```
kubeadm join 172.31.22.31:6443 --token gyakv9.hktjpt5usstl5u3y \
--discovery-token-ca-cert-hash sha256:2590fd7ba571e7e92b4f18f77c2149583f19f6049e3dfb4d306ac22cf2f465d6
```



```
[root@ip-172-31-23-217 ec2-user]# kubeadm join 172.31.22.31:6443 --token gyakv9.hktjpt5usstl5u3y \
--discovery-token-ca-cert-hash sha256:2590fd7ba571e7e92b4f18f77c2149583f19f6049e3dfb4d306ac22cf2f465d6
[preflight] Running pre-flight checks
```

Post which we are supposed to get the output that our worker nodes have been successfully connected to master node.

Unfortunately, on running the join command i was not able to produce anything beyond 'Running pre-flight checks' which can be seen in the above image

And thus, could not execute the last step of this experiment

**Conclusion:** In this experiment, we set up a connection between a local machine and an EC2 instance using SSH. After facing issues like timeouts and permission problems, we learned how to check for common causes such as incorrect security group settings, improper key permissions, and network issues. By resolving these, we successfully connected to the EC2 instance. This experiment helped us understand the steps required for remote server access and troubleshooting.